

# Decision-focused learning: theory, applications, challenges

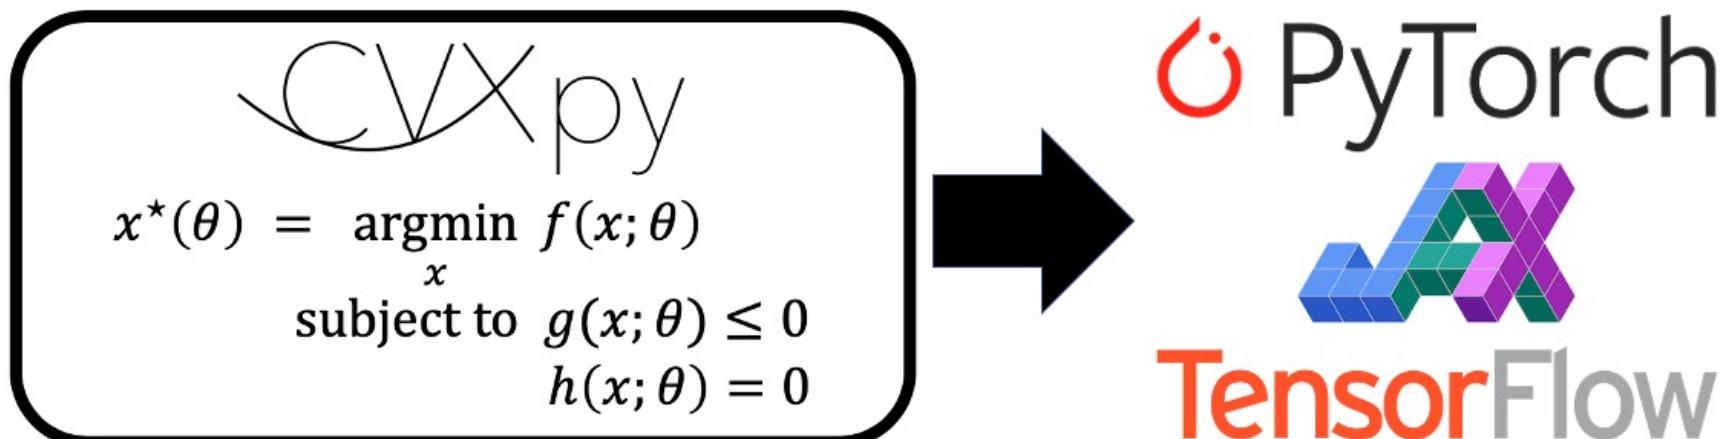
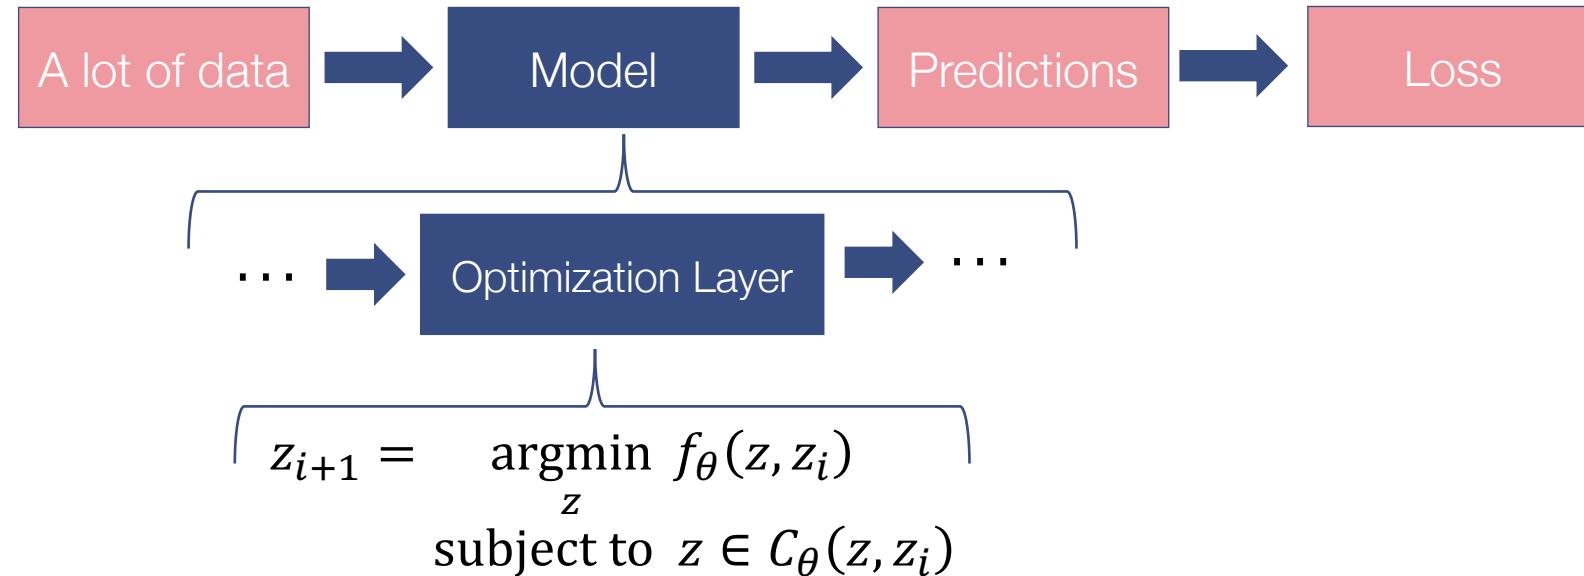
Andrew Perrault

Department of Computer Science and Engineering

The Ohio State University

IJCAI 2022

# Brandon's talk recap



# Agenda

What is decision-focused learning? Why should you care?

Decision-focused learning for sequential decision problems  
(NeurIPS 2021)

Learning convex & smooth loss functions for decision-focused learning (arXiv)

# Notation: predict-then-optimize

Observable feature/context vector  $x \in X$

Unobserved ground truth environment parameter  $\theta^* \in \Theta$

$x$  and  $\theta^*$  drawn jointly from fixed distribution  $\mathcal{D}$

Decision/control variable  $z \in Z$

Objective function  $f(z, \theta^*)$

Goal:  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

# Notation: predict-then-optimize

Observable feature/context vector  $x \in X$

Unobserved ground truth environment  
parameter  $\theta^* \in \Theta$

$x$  and  $\theta^*$  drawn jointly from fixed  
distribution  $\mathcal{D}$

Observed at training  
time, but not test time

Decision/control variable  $z \in Z$

Objective function  $f(z, \theta^*)$

Goal:  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

# Notation: predict-then-optimize

Observable feature/context vector  $x \in X$

Unobserved ground truth environment parameter  $\theta^* \in \Theta$

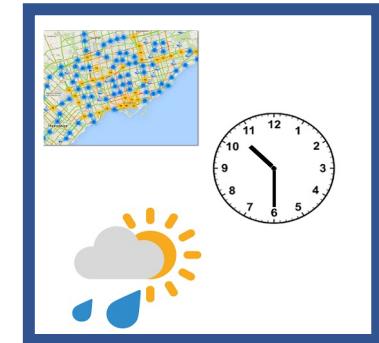
$x$  and  $\theta^*$  drawn jointly from fixed distribution  $\mathcal{D}$

Decision/control variable  $z \in Z$

Objective function  $f(z, \theta^*)$

Goal:  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

Context  $x$



Decision: route  $z$

Objective  $f(x, z)$ :  
minimize travel  
time under  
traffic

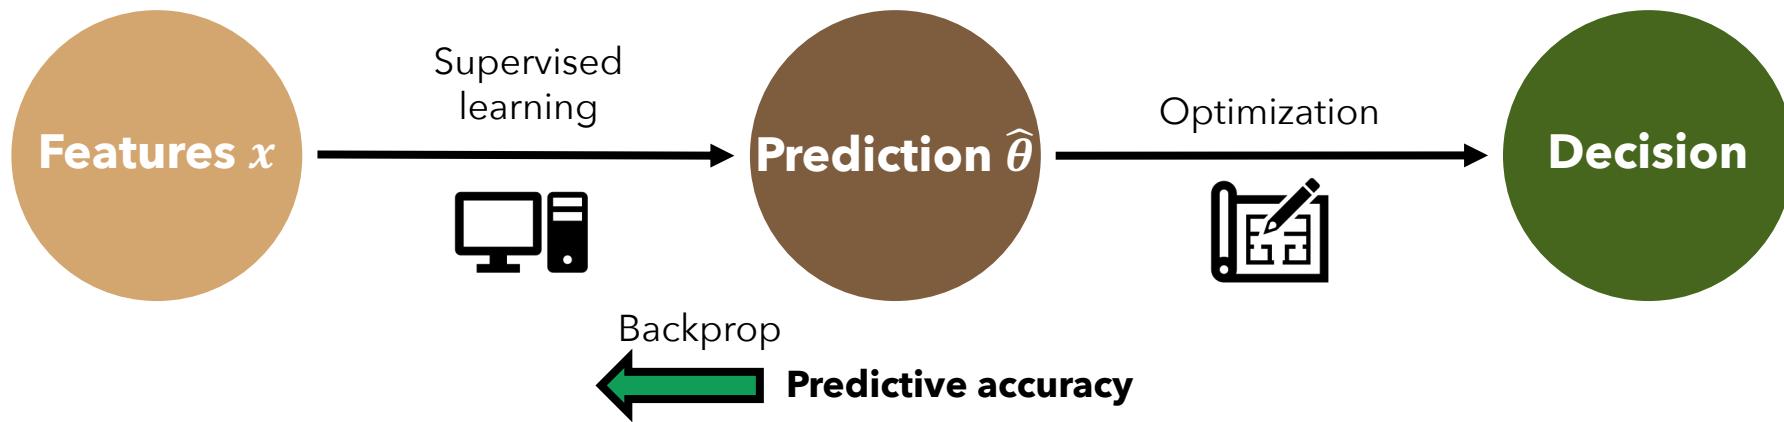


# Motivation of predict-then-optimize

Data =

{ (  i ,  i ) }  
i

# Two-stage pipeline for predict-then-optimize



# What do we lose by simplifying?

$\min_{z \in Z} f(z, \mathbb{E}(\theta|x))$  instead of  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

Two-stage

What we want

# What do we lose by simplifying?

$\min_{z \in Z} f(z, \mathbb{E}(\theta|x))$  instead of  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

Consider  $f(z, \theta) = \sum_i \theta_i z_i$ . Then,

$$\begin{aligned}\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta) &= \min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} \sum_i \theta_i z_i \\ &= \min_{z \in Z} \sum_i \mathbb{E}[\theta_i|x] z_i = \min_{z \in Z} f(z, \mathbb{E}(\theta|x))\end{aligned}$$

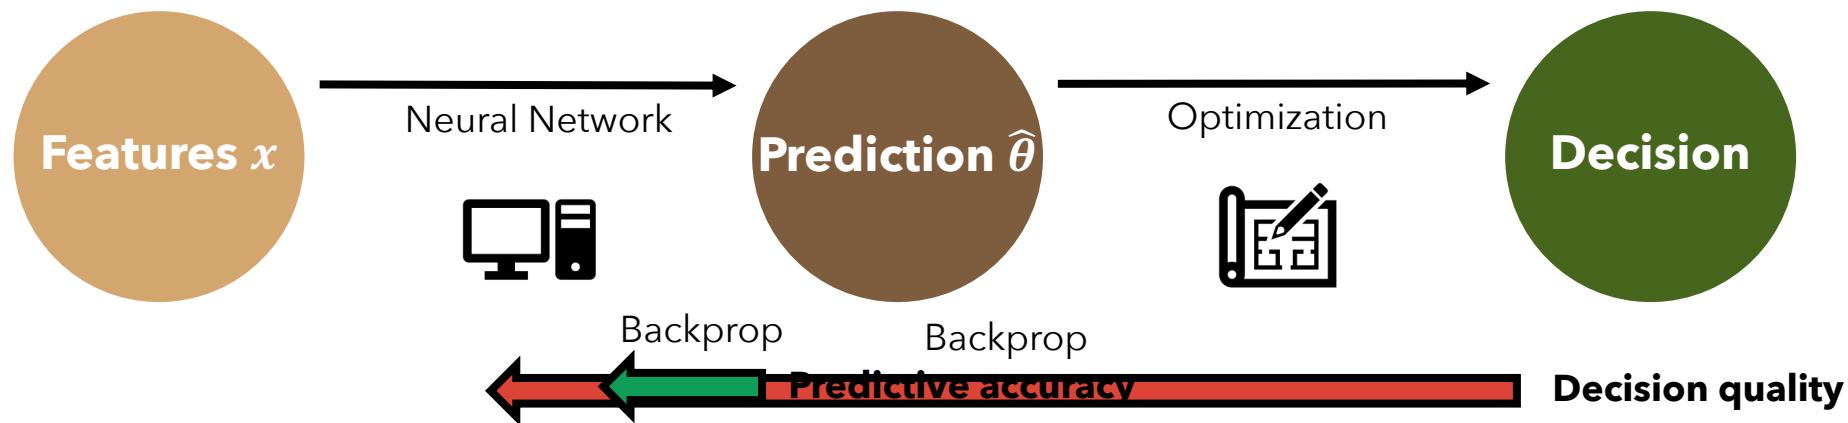
# Product of predictions

$\min_{z \in Z} f(z, \mathbb{E}(\theta|x))$  instead of  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

Consider  $f(z, \theta) = \sum_i \theta_{2i} \theta_{2i+1} z_{2i}$ . Then,

$$\begin{aligned} \min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta) &= \min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} \sum_i \theta_{2i} \theta_{2i+1} z_{2i} \\ &\neq \min_{z \in Z} \sum_i \mathbb{E}[\theta_{2i}|x] \mathbb{E}[\theta_{2i+1}|x] z_{2i} \end{aligned}$$

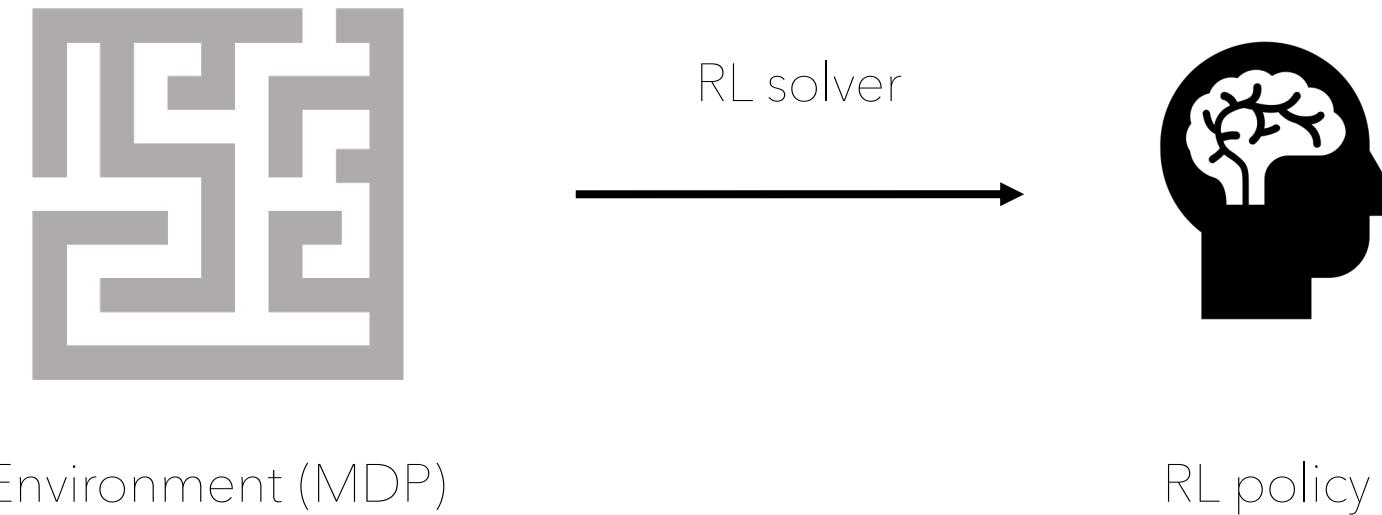
# Does differentiable optimization help?



Theorem (Cameron et al., 2022): There exists a class of problems (two-stage minimum cost flow from Agrawal et al., 2012) such that:

- $\text{DFL} = \text{OPT}$
- $\frac{2S}{\text{DFL}}$  is at least  $\Omega(2^d)$ , where  $d = \dim(Z) = \dim(\Theta)$

# Optimizing policies in unknown MDPs



# Planning for Community Health Workers

Key tasks:

Maternal health checkups

Behavioral health interventions

Medication adherence  
support



Photo credit: Pippa Ranger

# Adherence monitoring for TB

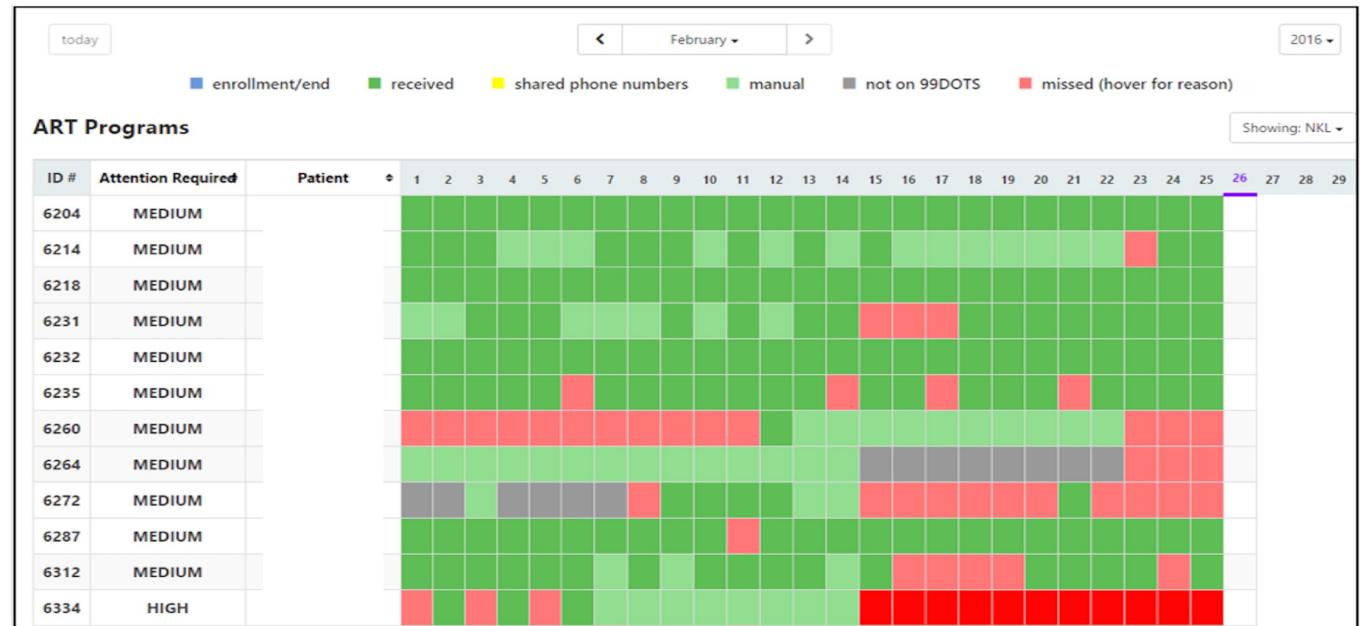
Tuberculosis (TB): 500,000 deaths/year, 3 million infected in India

Curing TB requires 6 months of daily medication

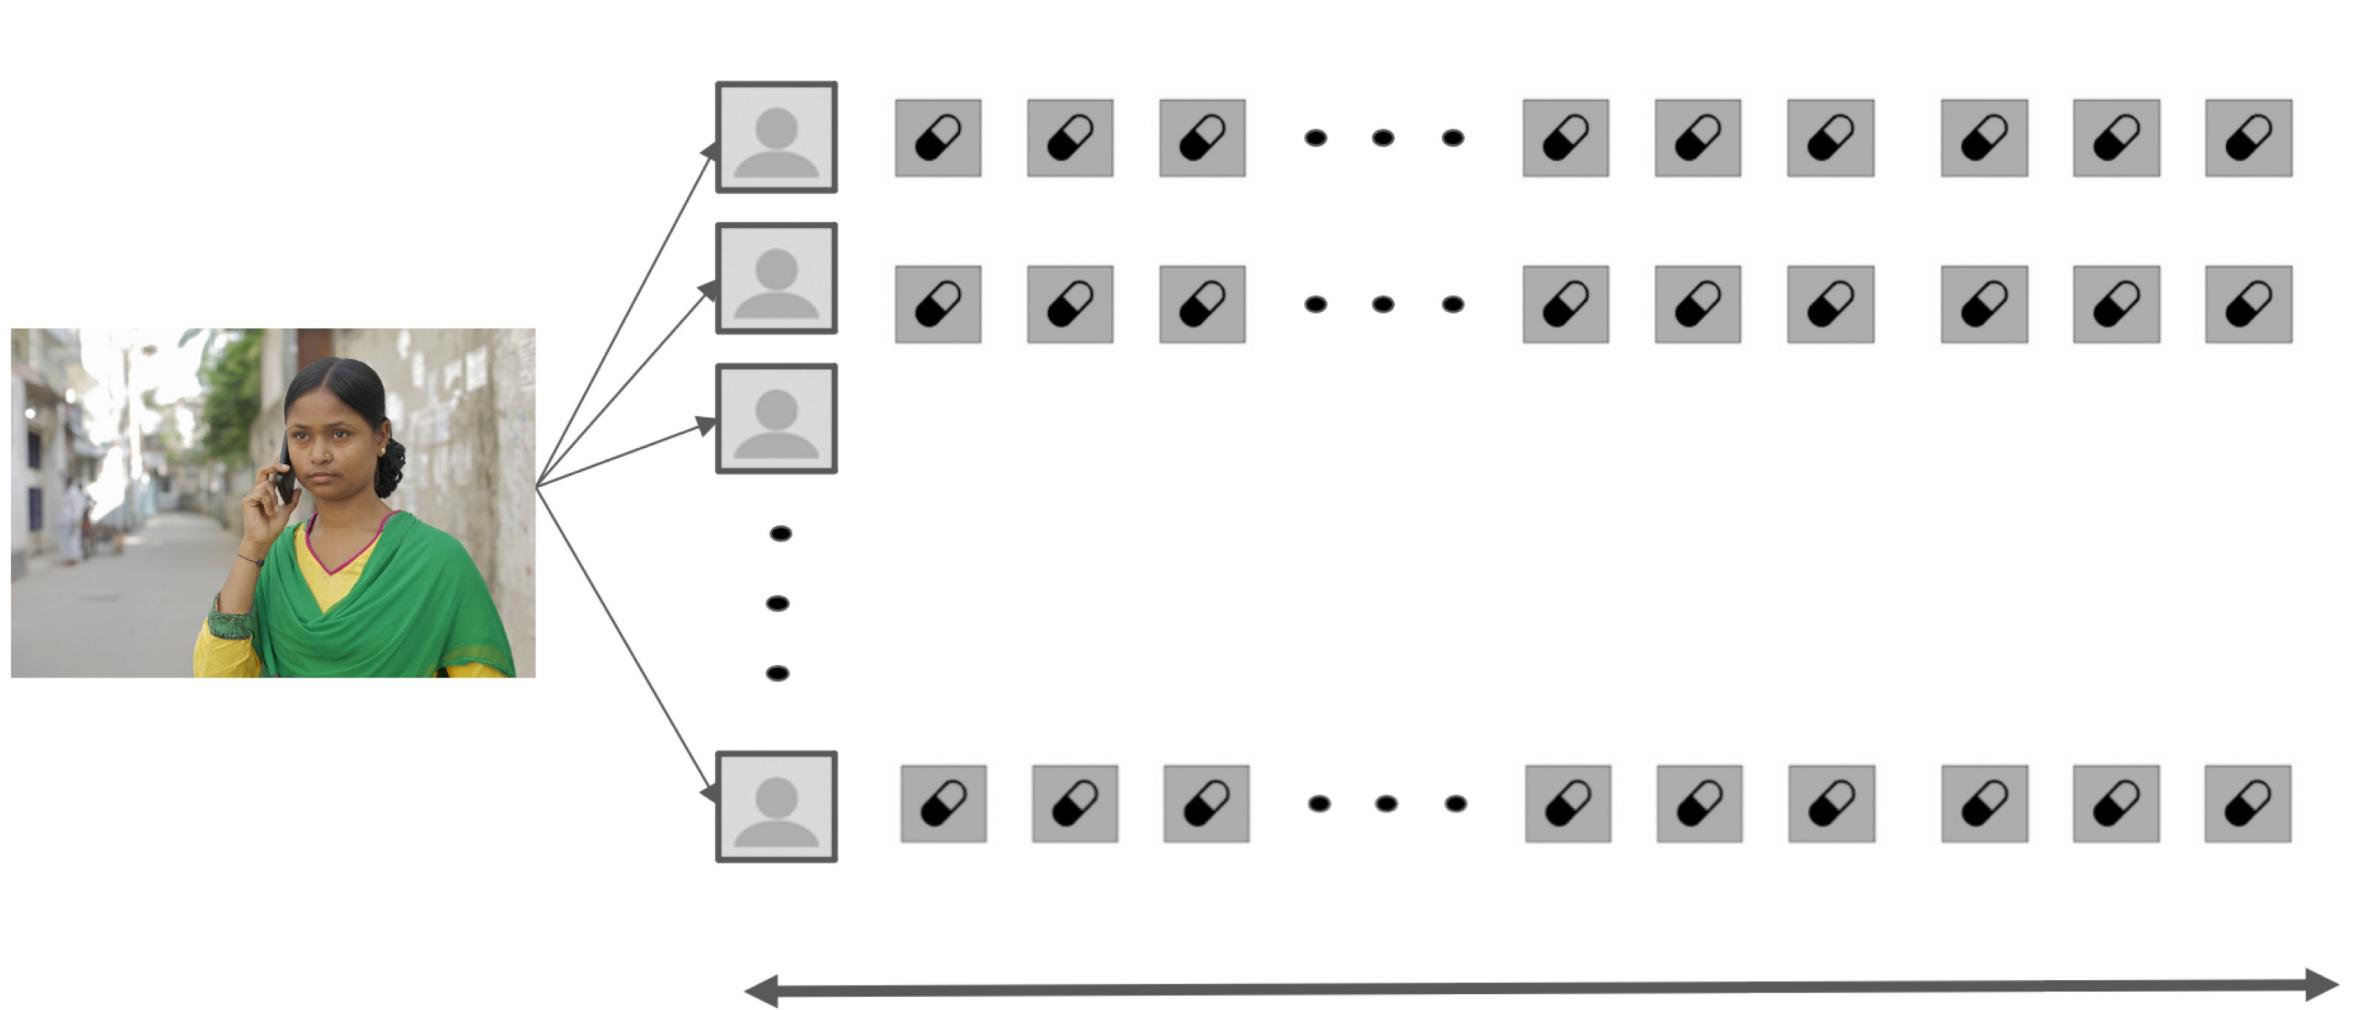
Symptoms disappear quickly



# TB medication adherence

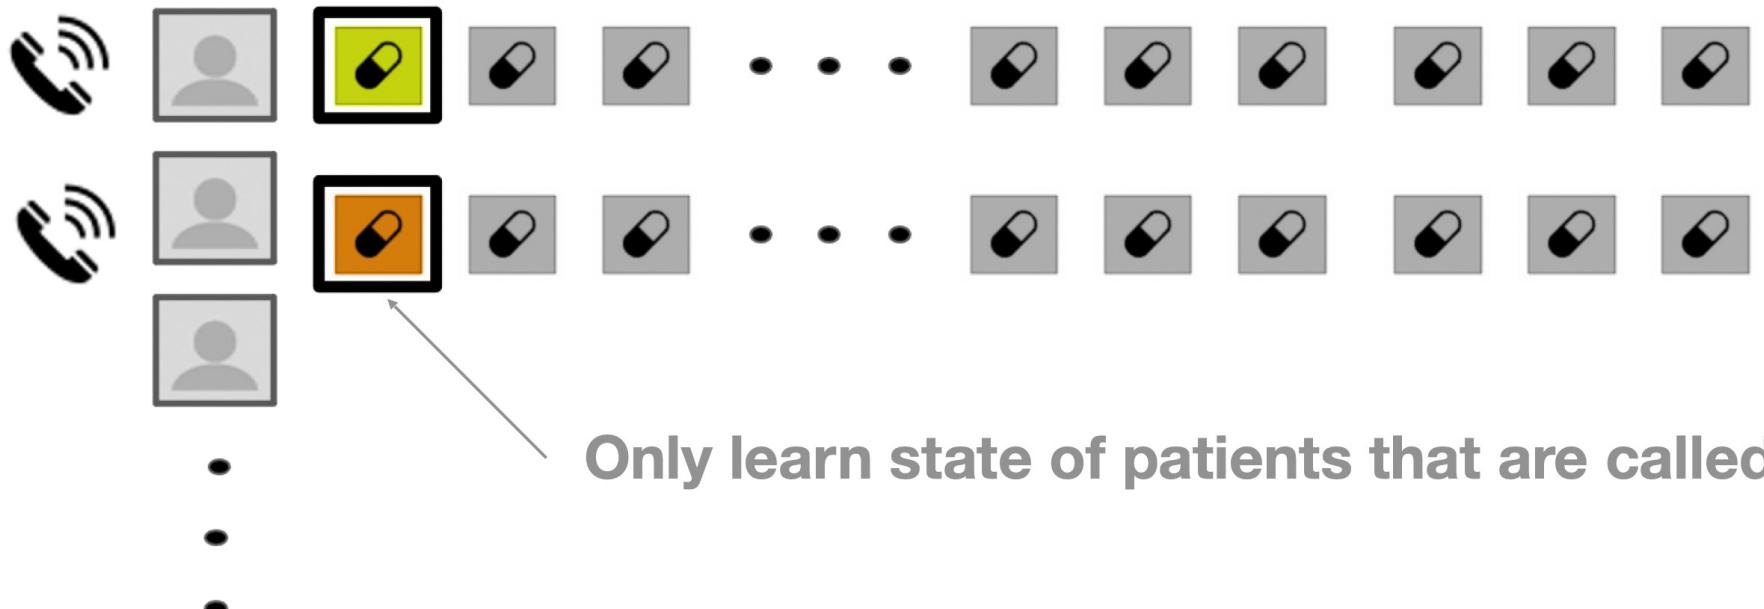


# Adherence example



**6 months**

# Adherence example



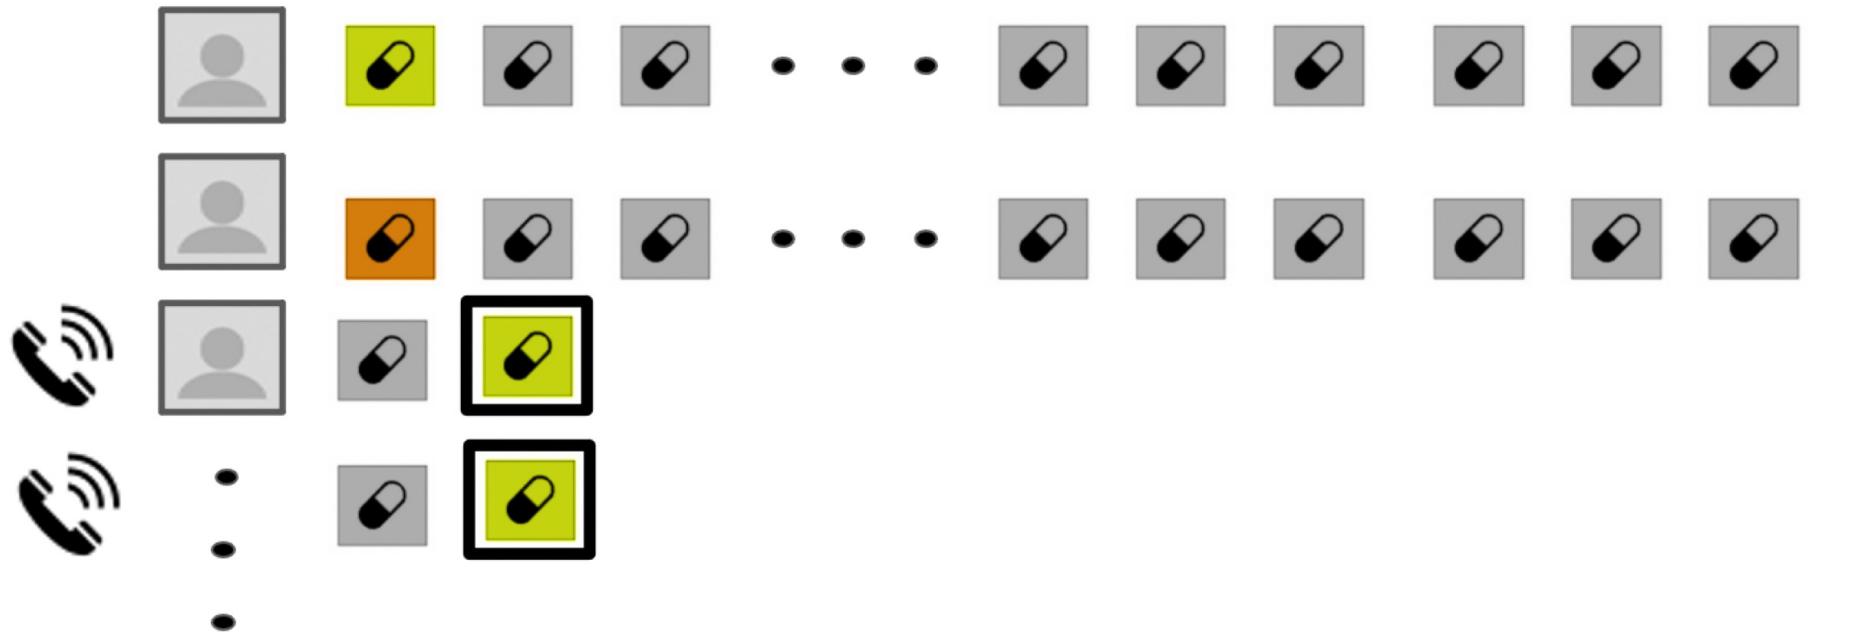
**Adhering**

**Not adhering**

**Unknown**



# Adherence example



Adhering



Not adhering



Unknown



# POMDP formulation

Each patient is represented by a partially observable Markov decision process (POMDP)

Shared action budget across all processes

States



Actions



Observations



Transitions

$$P_n^{\text{call}} \quad P_n^{\text{no call}}$$

Rewards

$$\begin{array}{ll} \text{yellow capsule} = 1 & \text{orange capsule} = 0 \end{array}$$

Goal: find optimal policy  $\pi^*$ —a mapping from the action-observation history to actions that maximizes reward

# POMDP formulation

Each patient is represented by a partially observable Markov decision process (POMDP)

Shared action budget across all processes

States



Actions



Observations



Unknown and different per patient!

Transitions



Goal: find optimal policy  $\pi^*$ —a mapping from the action-observation history to actions that maximizes reward

Rewards



# Patient data

Patient context  $x$

Trajectory  $\{(s_t, a_t, r_t, s_{t+1})\}$



# Predict-then-optimize?

Observable feature/context vector

$$x \in X$$

Unobserved ground truth

environment parameter  $\theta^* \in \Theta$

$x$  and  $\theta^*$  drawn jointly from fixed  
distribution  $\mathcal{D}$

Decision/control variable  $z \in Z$

Objective function  $f(z, \theta^*)$

Goal:  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

# Predict-then-optimize?

Observable feature/context vector

$$x \in X$$

Unobserved ground truth

environment parameter

$$\theta^* \in \Theta$$

$x$  and  $\theta^*$  drawn jointly from fixed  
distribution  $\mathcal{D}$

Not available at training time—  
only have trajectories

Decision/control variable  $z \in Z$

Objective function  $f(z, \theta^*)$

Goal:  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

# Contributions

How to differentiate through the optimal policy in sequential decision problems  
(and how to speed it up)

Do DFL with only trajectories (i.e., without  $\theta^*$  labels)

# Problem setting

Let  $(\mathcal{S}, s_0, A, T, R)$  be an *incomplete* MDP specification with:

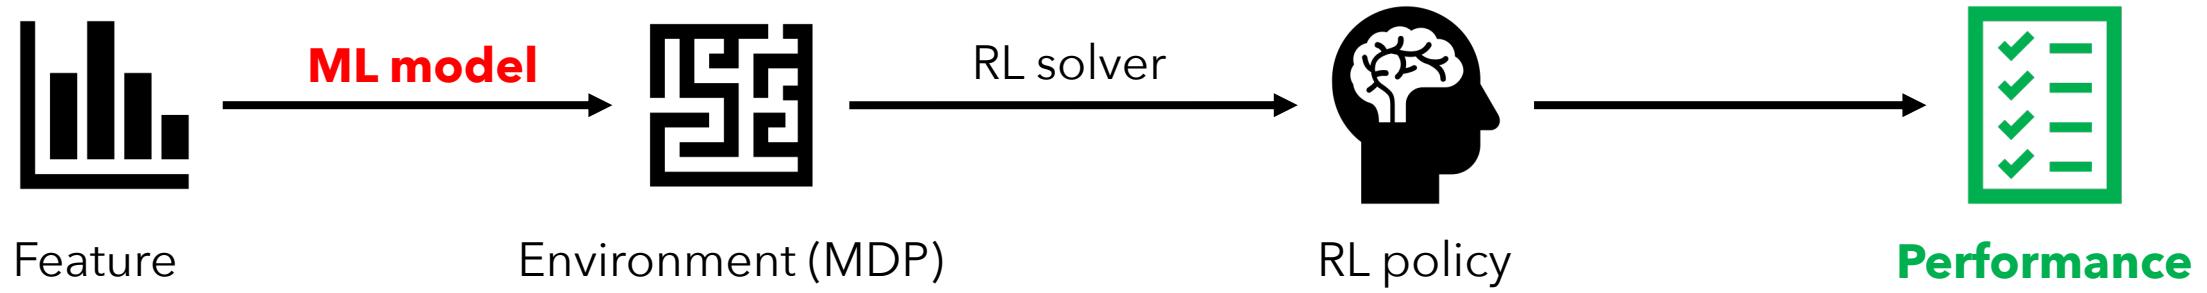
- Initial state  $s_0$
- Set of states  $\mathcal{S}$
- Set of actions  $A$
- Possibly incomplete transition function:  $T: \mathcal{S} \times A \rightarrow \Delta \mathcal{S}$
- Possible incomplete reward function:  $R: \mathcal{S} \times A \rightarrow \mathbb{R}$

Each MDP has a context vector  $x$

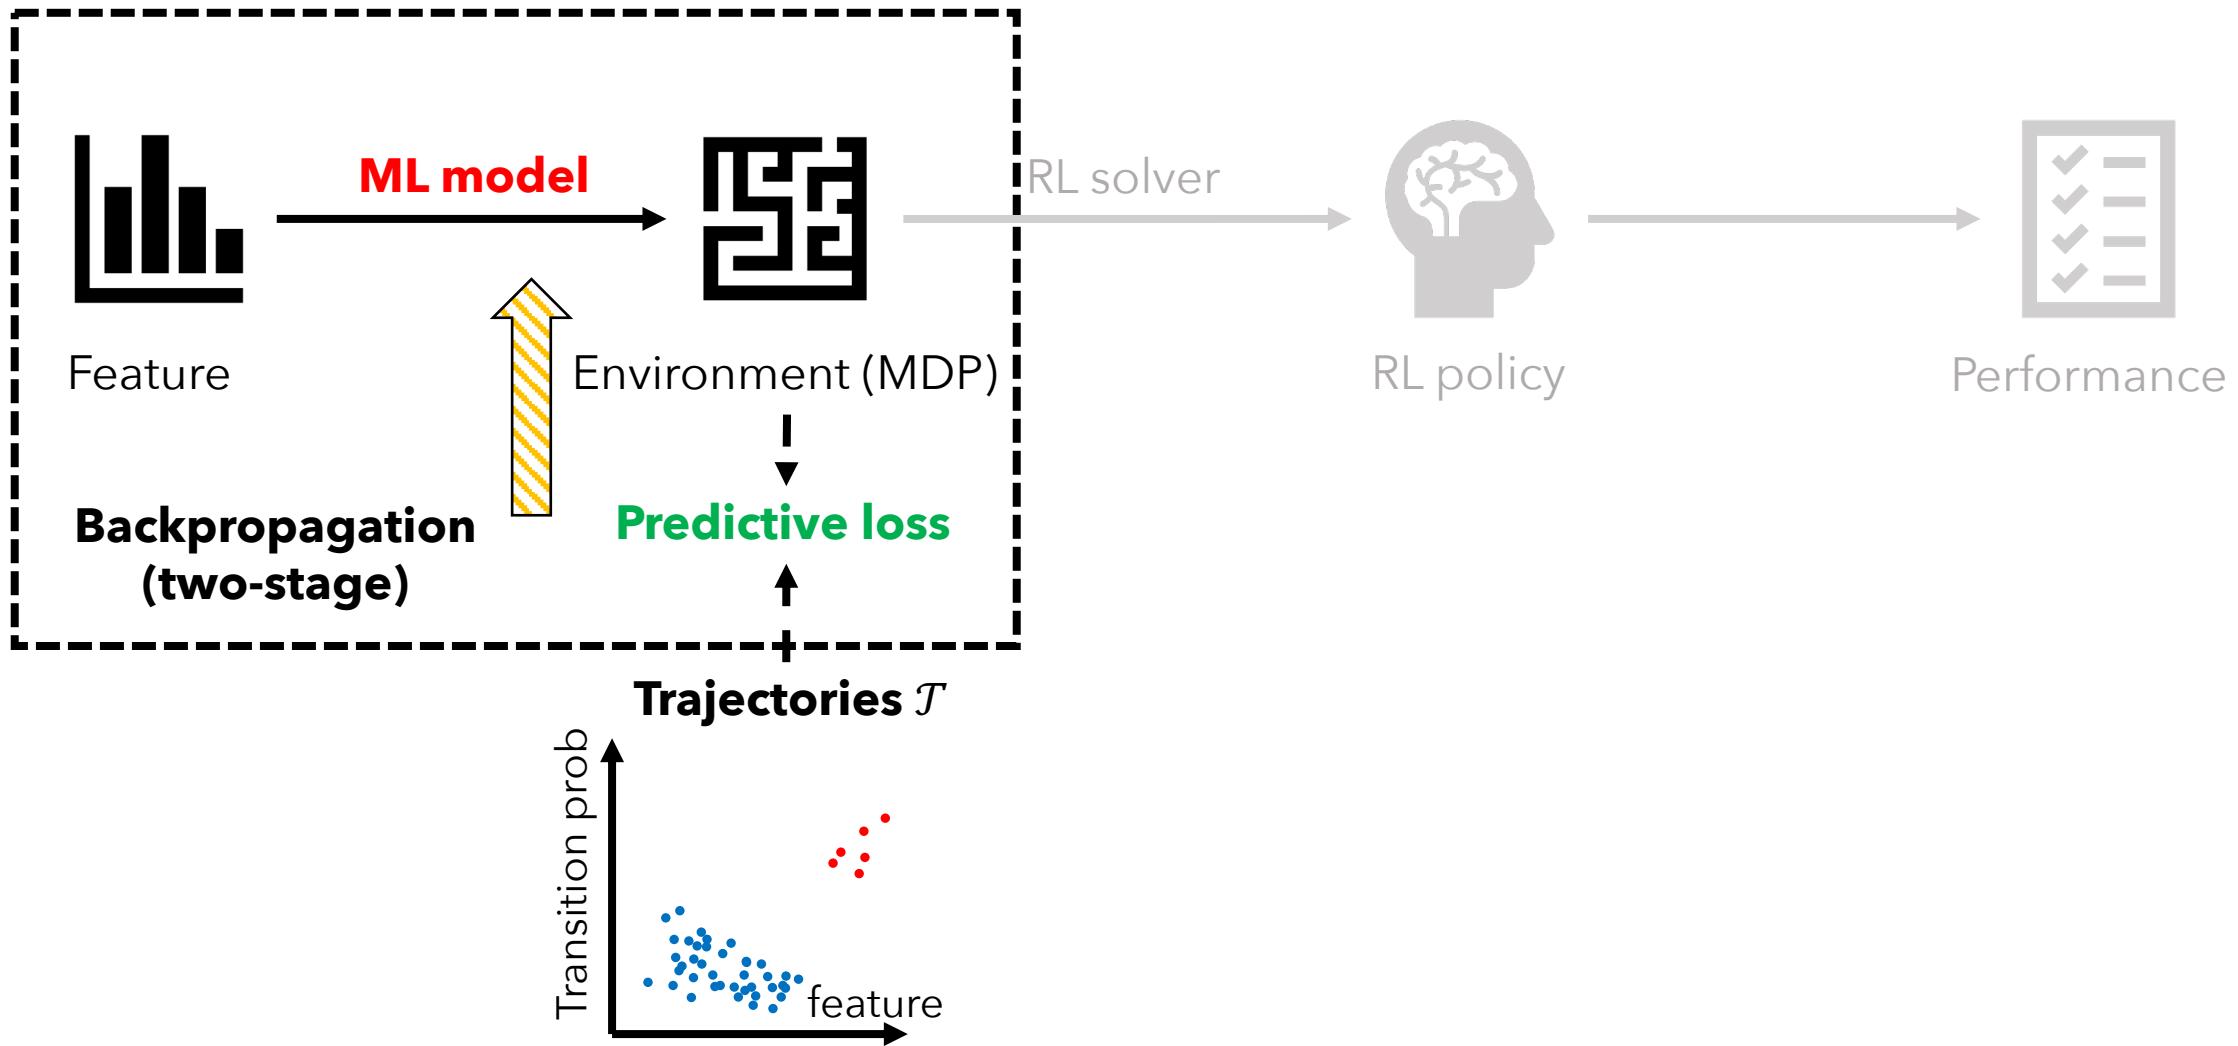
Each training MDP has a set of trajectories  $\mathcal{T}$

Goal: learn a predictive model  $m_w$  that “completes” the MDP from  $x$  and maximizes reward on  $\mathcal{T}$

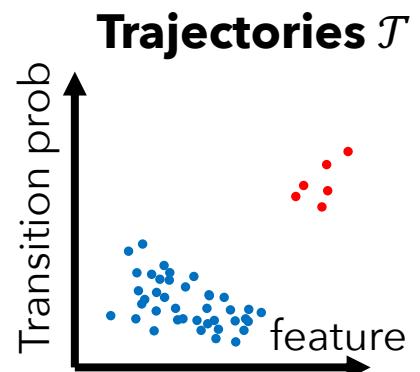
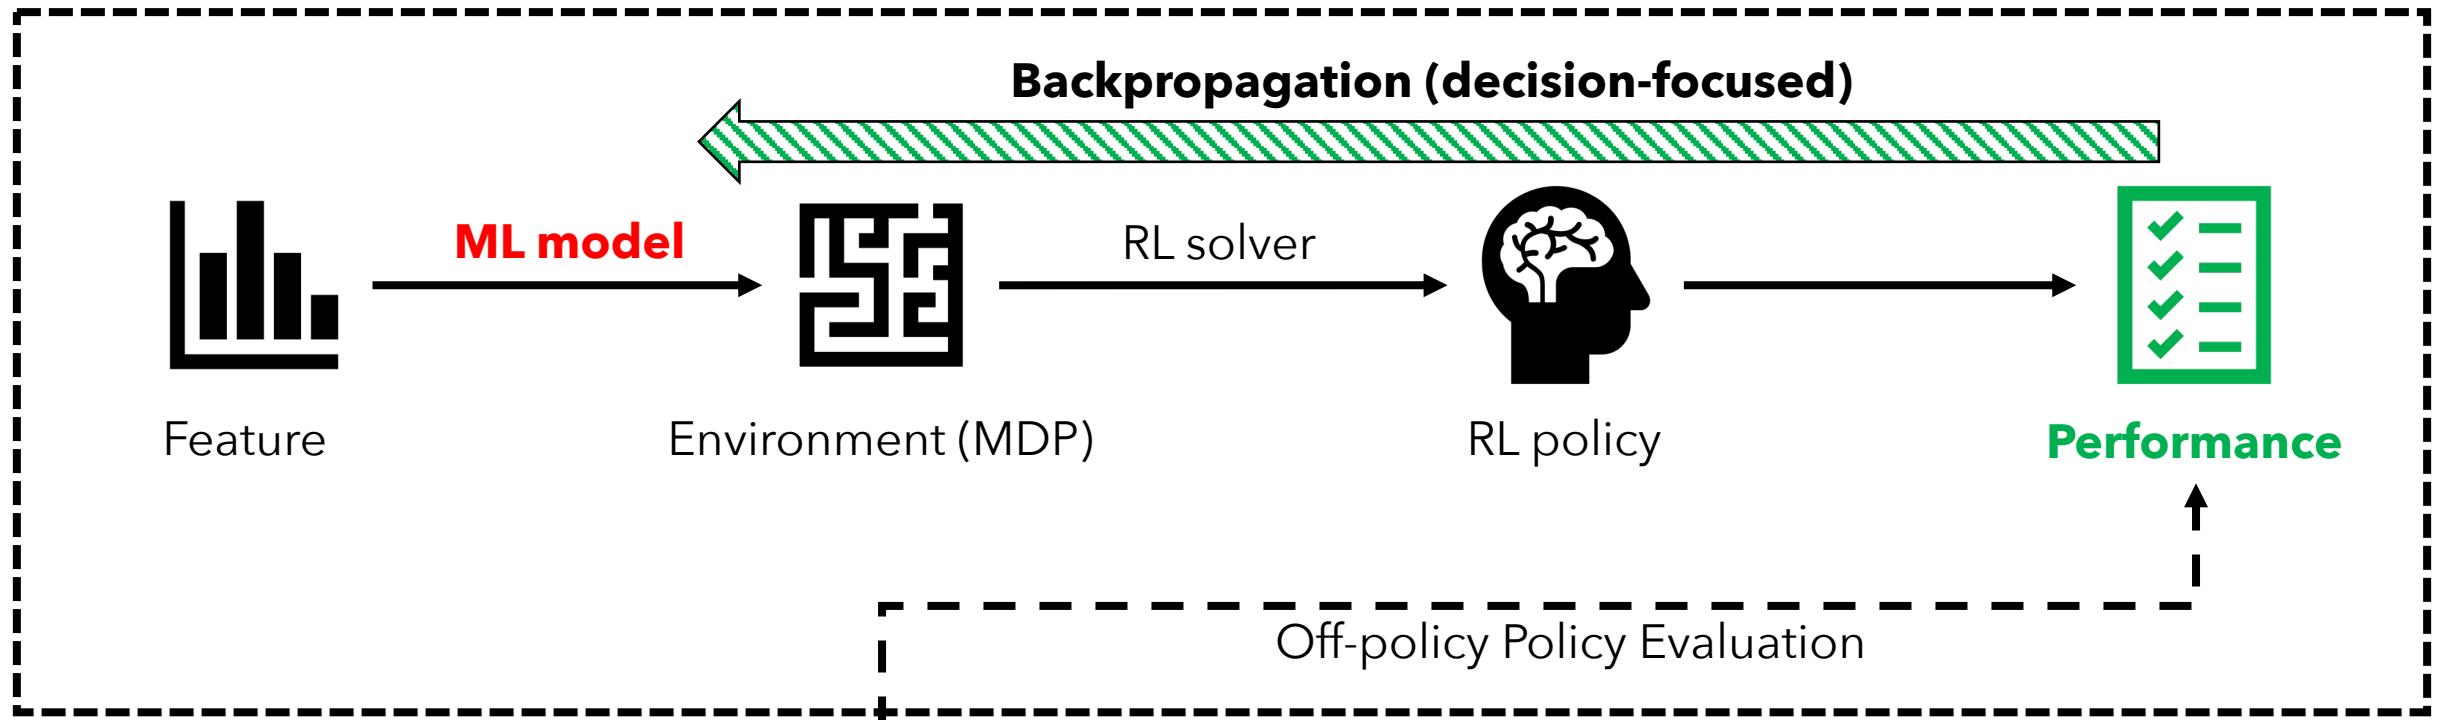
# Pipeline visualized



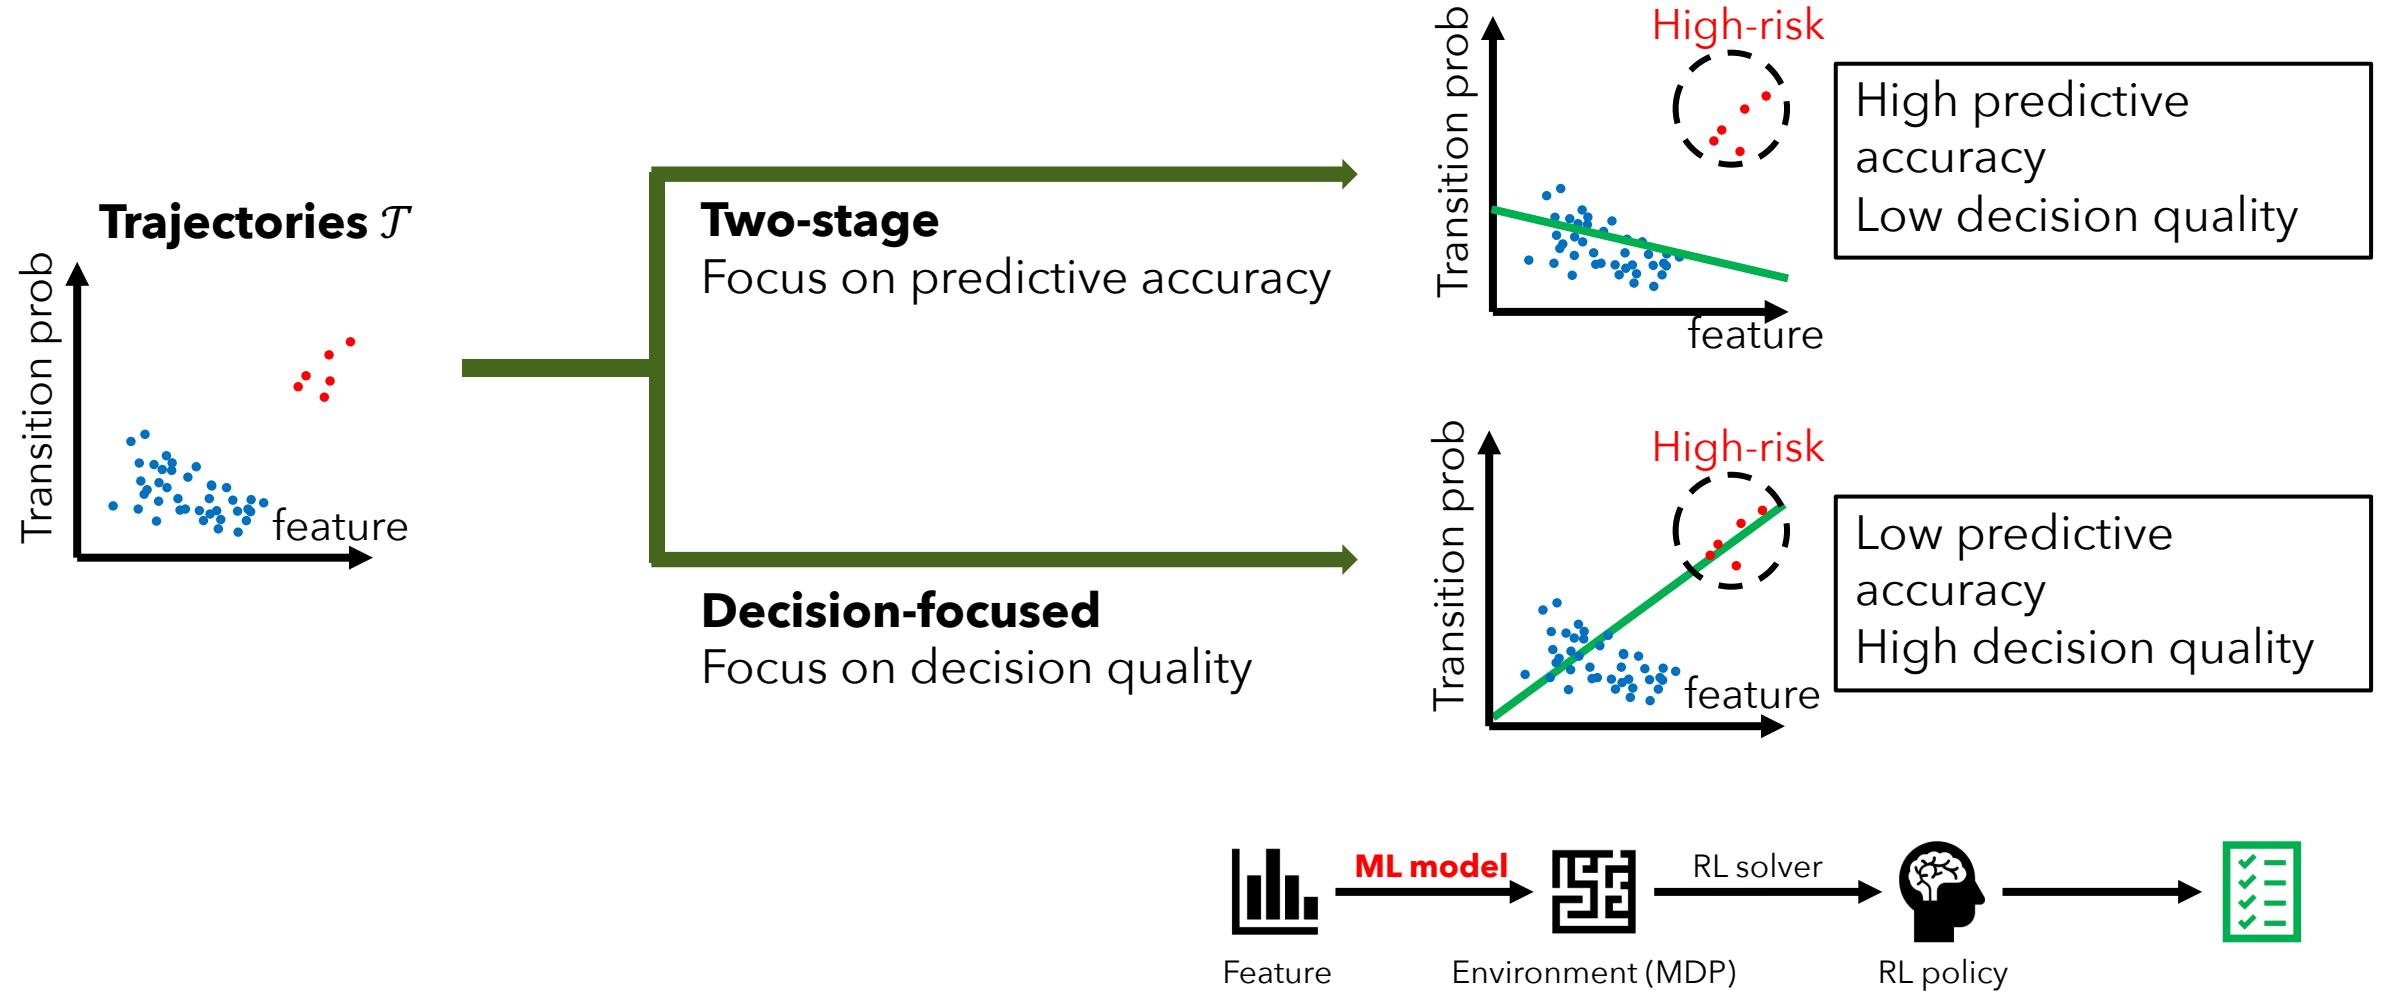
# Two-stage approach



# Decision-focused approach



# Two-stage vs. DFL

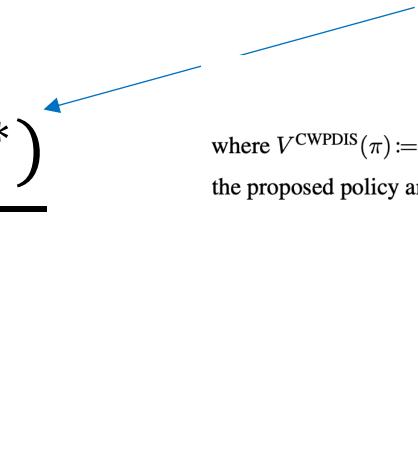


# How DFL works

To run gradient descent, we need to compute

$$\frac{d\text{Eval}_T(\pi^*)}{dw} = \frac{d\theta}{dw} \frac{d\pi^*}{d\theta} \frac{d\text{Eval}_T(\pi^*)}{d\pi^*}$$

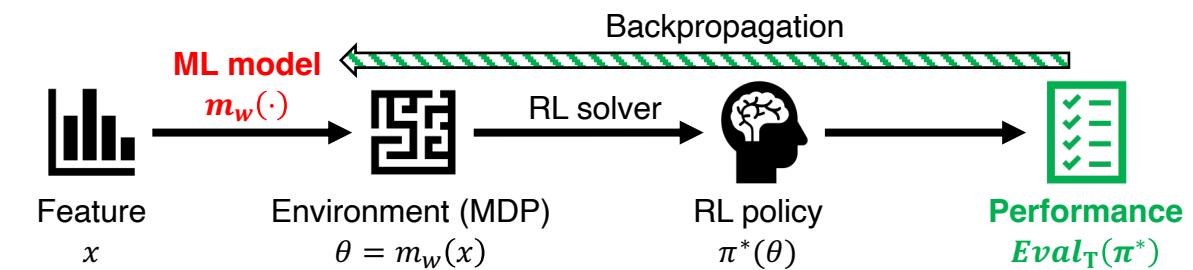
Weights  $w$  of  
predictive model



where  $V^{\text{CWPDIS}}(\pi) := \sum_{t=1}^h \gamma^t \frac{\sum_i r_{it} \rho_{it}(\pi)}{\sum_i \rho_{it}(\pi)}$  and  $\text{ESS}(\pi) := \sum_{t=1}^h \frac{(\sum_i \rho_{it})^2}{\sum_i \rho_{it}^2}$ , and  $\rho_{it}(\pi)$  is the ratio of the proposed policy and the behavior policy likelihoods up to time  $t$ :  $\rho_{it}(\pi) := \prod_{t'=1}^t \frac{\pi(a_{it'}|s_{it'})}{\pi_{\text{beh}}(a_{it'}|s_{it'})}$ .

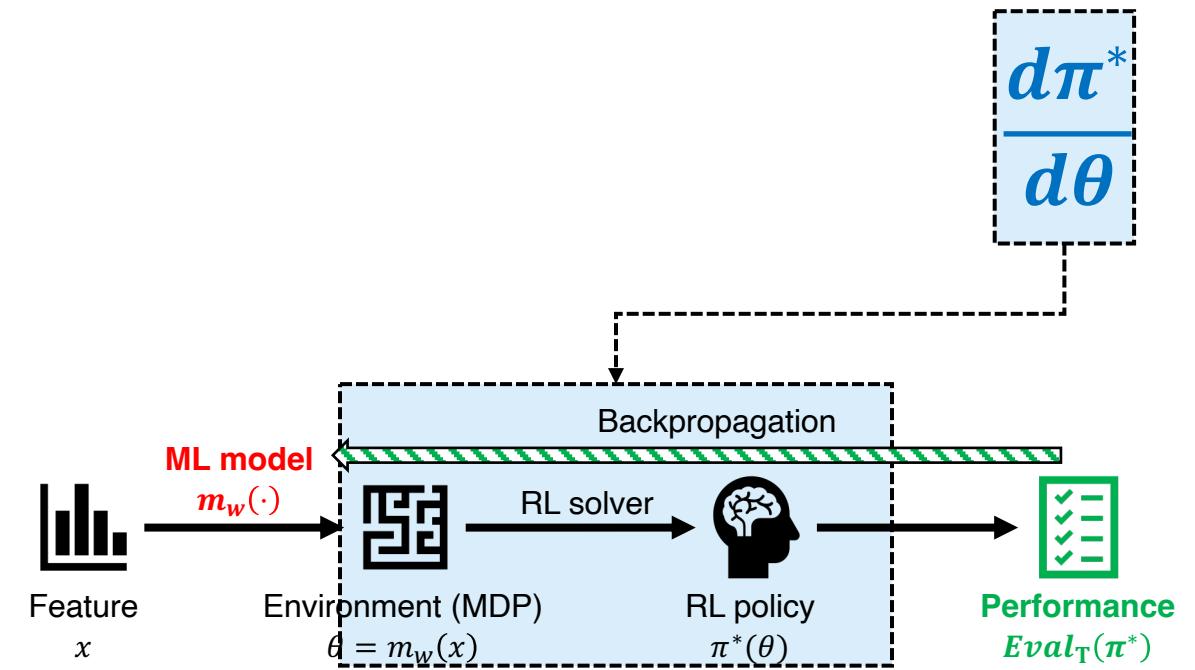
We use a metric from  
Futoma et al.

$$\text{Eval}_T(\pi) := V^{\text{CWPDIS}}(\pi) - \lambda_{\text{ESS}} / \sqrt{\text{ESS}(\pi)} \quad (1)$$



# Leveraging optimality conditions

Differentiate through one-shot optimization problem by optimality condition (Amos et al., 2017)



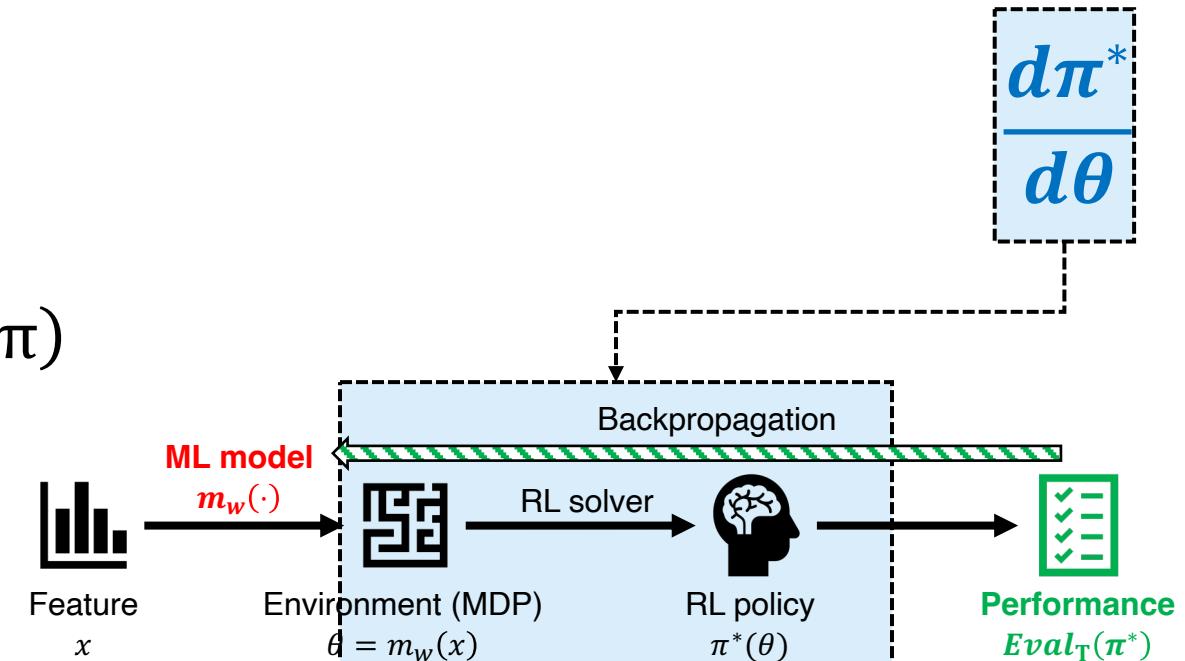
# Leveraging optimality conditions

Policy optimality:  $\pi^* = \operatorname{argmax}_{\pi} J_{\theta}(\pi)$

Expected cumulative return  
↓

Bellman optimality:  $\pi^* = \operatorname{argmin}_{\pi} B_{\theta}(\pi)$

Differentiate through one-shot optimization problem  
by optimality condition (Amos et al., 2017)



# Leveraging optimality conditions

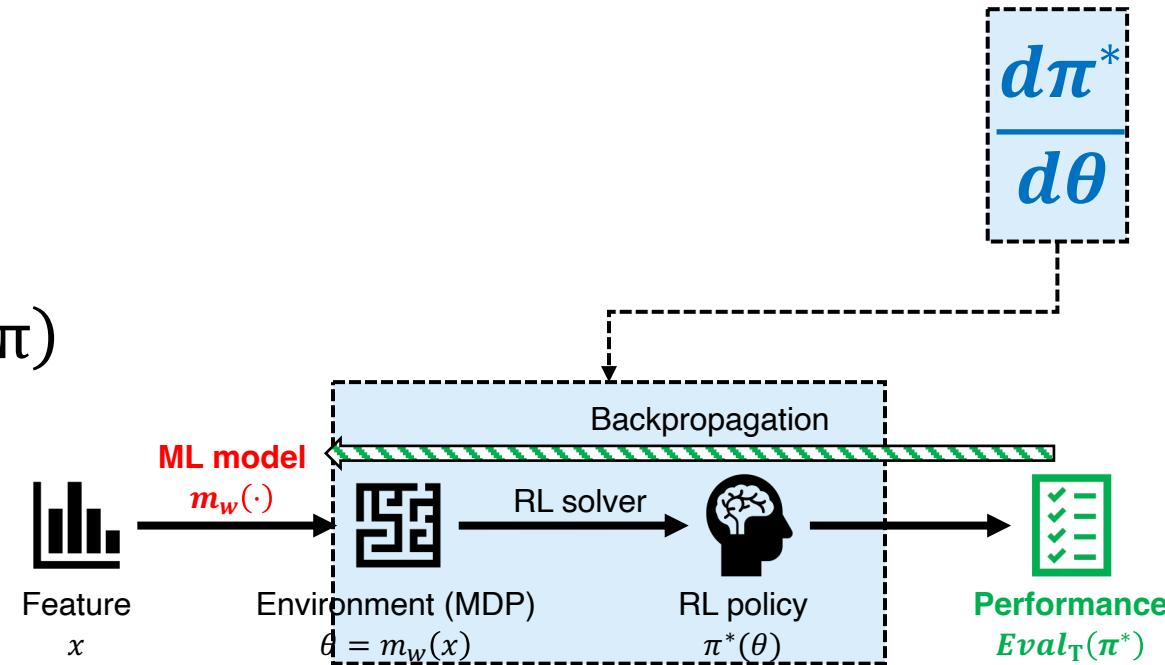
Policy optimality:  $\pi^* = \operatorname{argmax}_{\pi} J_{\theta}(\pi)$

KKT conditions:  $\nabla_{\pi} J_{\theta}(\pi^*) = 0$

Bellman optimality:  $\pi^* = \operatorname{argmin}_{\pi} B_{\theta}(\pi)$

Differentiate through one-shot optimization problem by optimality condition (Amos et al., 2017)

Expected cumulative return  
↓



# Leveraging optimality conditions

Policy optimality:  $\pi^* = \operatorname{argmax}_{\pi} J_{\theta}(\pi)$

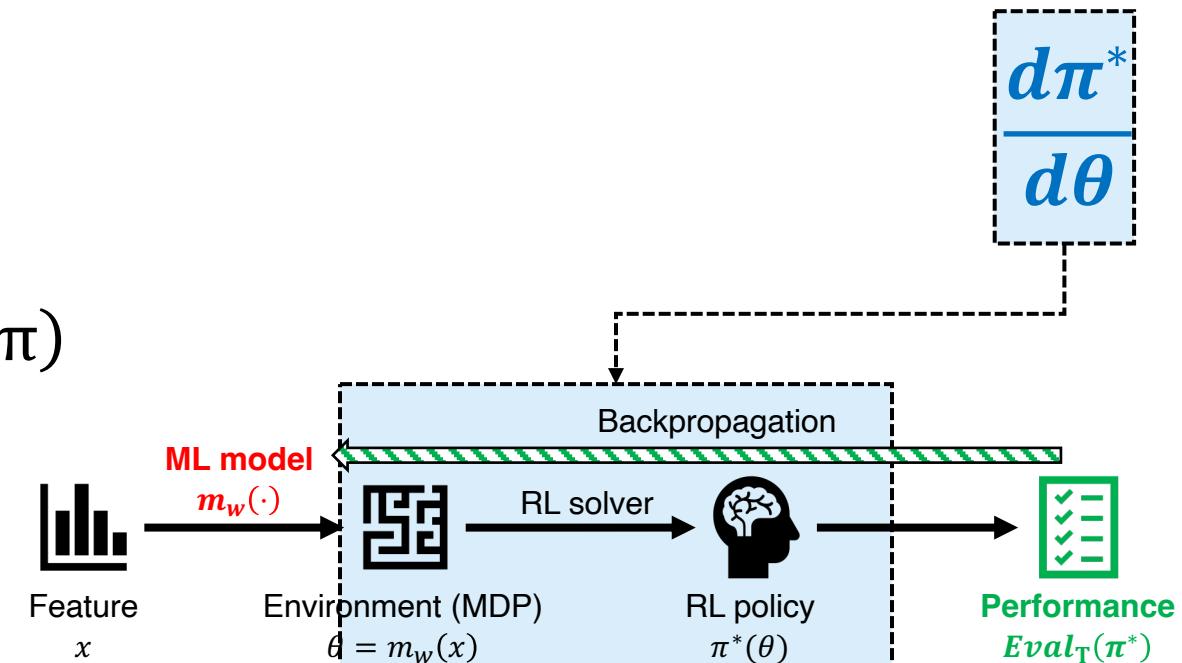
KKT conditions:  $\frac{d}{d\theta} \nabla_{\pi} J_{\theta}(\pi^*) = 0$

Bellman optimality:  $\pi^* = \operatorname{argmin}_{\pi} B_{\theta}(\pi)$

Differentiate through one-shot optimization problem by optimality condition (Amos et al., 2017)

Expected cumulative return  
↓

$$\frac{d\pi^*}{d\theta}$$



# Leveraging optimality conditions

Policy optimality:  $\pi^* = \operatorname{argmax}_{\pi} J_{\theta}(\pi)$

KKT conditions:  $\frac{d}{d\theta} \nabla_{\pi} J_{\theta}(\pi^*) = 0$

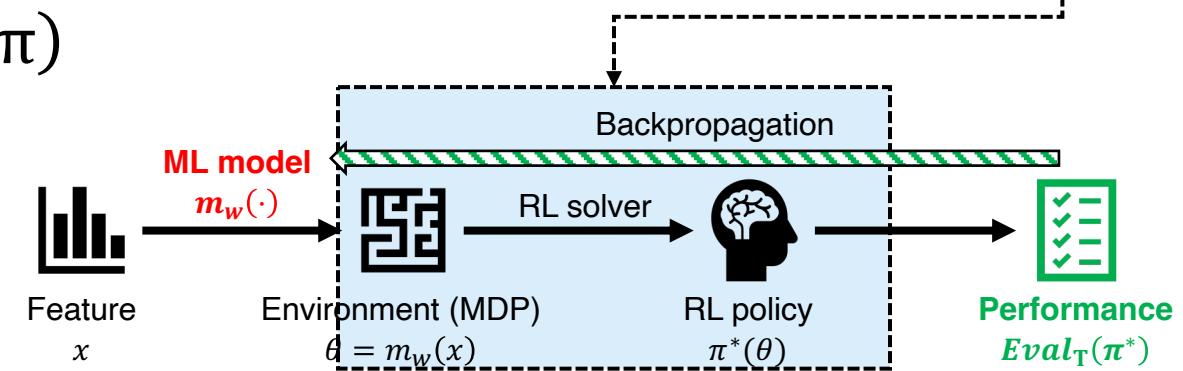
$$\nabla_{\pi}^2 J_{\theta}(\pi^*) \frac{d\pi^*}{d\theta} + \nabla_{\theta\pi}^2 J_{\theta}(\pi^*) = 0$$

Bellman optimality:  $\pi^* = \operatorname{argmin}_{\pi} B_{\theta}(\pi)$

Differentiate through one-shot optimization problem by optimality condition (Amos et al., 2017)

Expected cumulative return  
↓

$$\boxed{\frac{d\pi^*}{d\theta}}$$



# Leveraging optimality conditions

Policy optimality:  $\pi^* = \operatorname{argmax}_{\pi} J_{\theta}(\pi)$

KKT conditions:  $\frac{d}{d\theta} \nabla_{\pi} J_{\theta}(\pi^*) = 0$

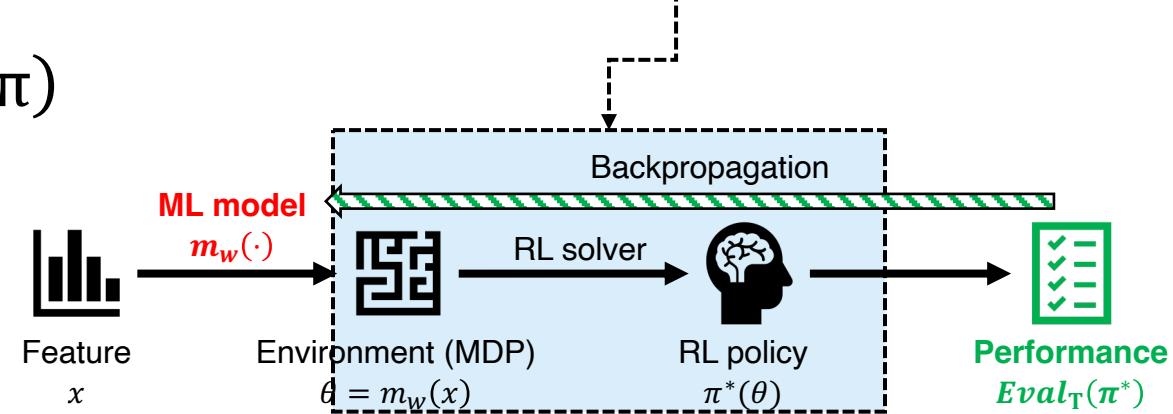
$$\nabla_{\pi}^2 J_{\theta}(\pi^*) \frac{d\pi^*}{d\theta} + \nabla_{\theta\pi}^2 J_{\theta}(\pi^*) = 0$$

Bellman optimality:  $\pi^* = \operatorname{argmin}_{\pi} B_{\theta}(\pi)$

Differentiate through one-shot optimization problem by optimality condition (Amos et al., 2017)

Expected cumulative return (implicitly defined)

$$\frac{d\pi^*}{d\theta} = (\nabla_{\pi}^2 J_{\theta}(\pi^*))^{-1} \nabla_{\theta\pi}^2 J_{\theta}(\pi^*)$$



# Computational challenges

Policy optimality:  $\pi^* = \operatorname{argmax}_{\pi} J_{\theta}(\pi)$

Bellman optimality:  $\pi^* = \operatorname{argmin}_{\pi} B_{\theta}(\pi)$

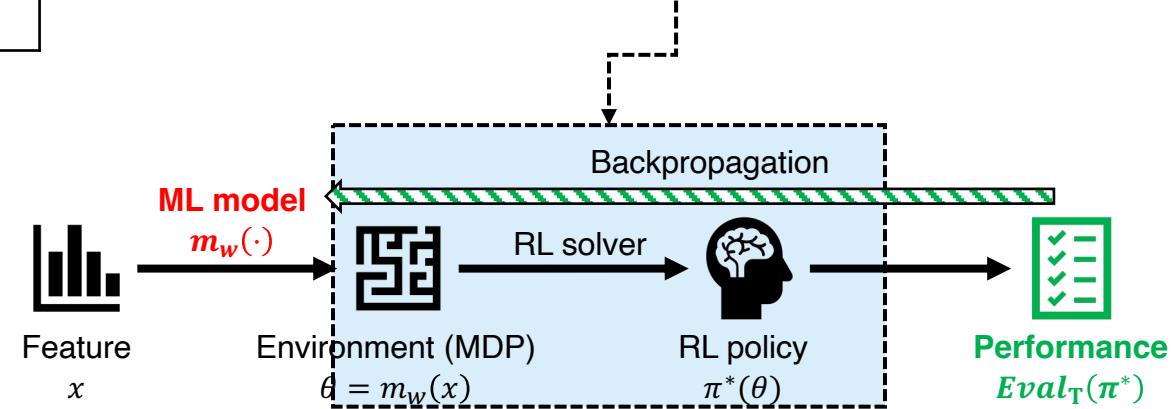
Different approximations of  $\nabla_{\pi}^2 J_{\theta}(\pi^*)$

Method	Computation	Accuracy
Full Hessian	$n^2 k + n^{\omega}$ <span style="color:red">✗</span>	High

$n$ : size of  $\pi$

$k$ : number of trajectories to compute the policy gradient

$$\frac{d\pi^*}{d\theta} = (\nabla_{\pi}^2 J_{\theta}(\pi^*))^{-1} \nabla_{\theta\pi}^2 J_{\theta}(\pi^*)$$



# Computational challenges

Policy optimality:  $\pi^* = \operatorname{argmax}_\pi J_\theta(\pi)$

Bellman optimality:  $\pi^* = \operatorname{argmin}_\pi B_\theta(\pi)$

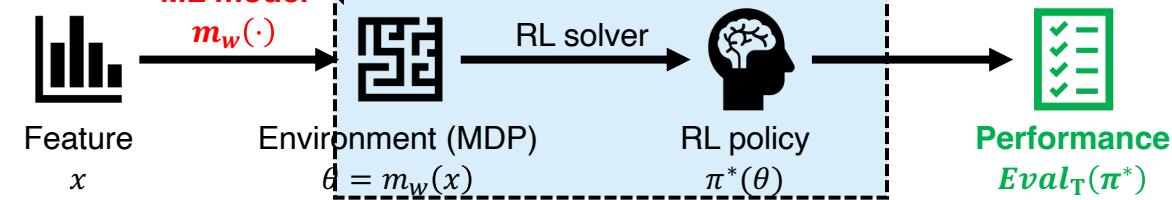
Different approximations of  $\nabla_\pi^2 J_\theta(\pi^*)$

Method	Computation	Accuracy
Full Hessian	$n^2 k + n^\omega \times$	High
Identity matrix	$nk \checkmark$	Low

$n$ : size of  $\pi$

$k$ : number of trajectories to compute the policy gradient

$$\frac{d\pi^*}{d\theta} = (\nabla_\pi^2 J_\theta(\pi^*))^{-1} \nabla_{\theta\pi}^2 J_\theta(\pi^*)$$



# Computational challenges

Policy optimality:  $\pi^* = \operatorname{argmax}_\pi J_\theta(\pi)$

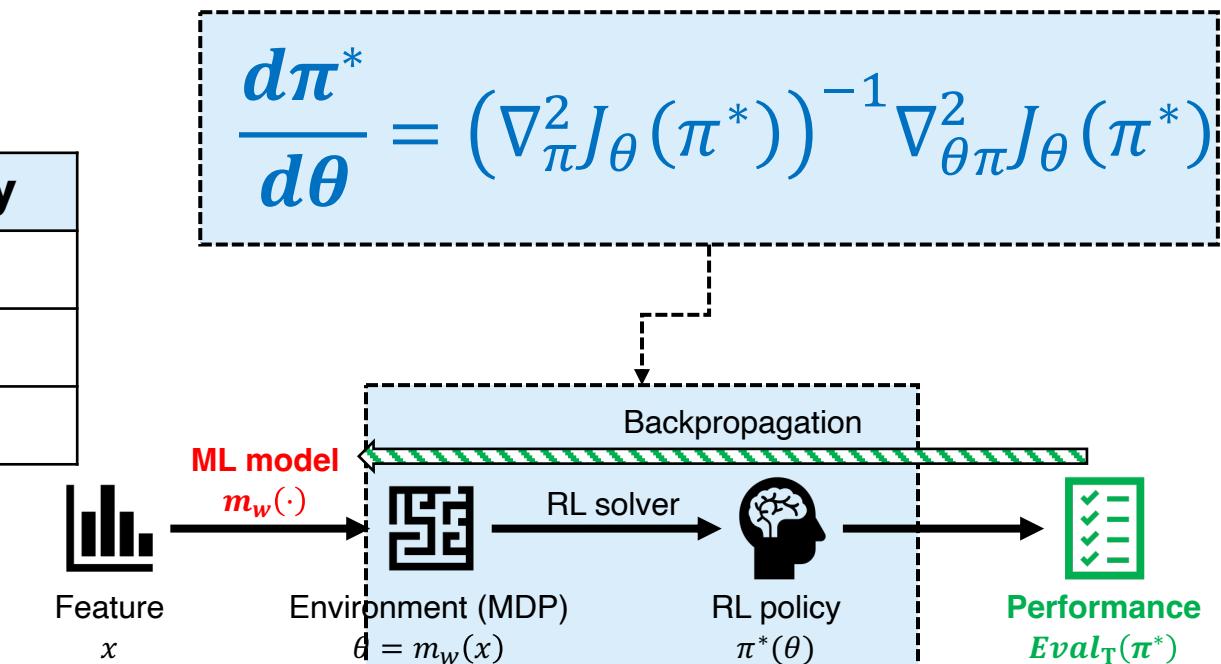
Bellman optimality:  $\pi^* = \operatorname{argmin}_\pi B_\theta(\pi)$

Different approximations of  $\nabla_\pi^2 J_\theta(\pi^*)$

Method	Computation	Accuracy
Full Hessian	$n^2 k + n^\omega \times$	High
Identity matrix	$nk \quad \checkmark$	Low
Woodbury (low-rank $k$ )	$nk + k^\omega \checkmark$	Medium

$n$ : size of  $\pi$

$k$ : number of trajectories to compute the policy gradient



# Experiments

Predict-then-optimize problems with missing MDP parameters

Reward

0.5	0	0	1	10
0.5	-8	0.5	1	-3
2	0.5	0.5	1	-10
0.5	-10	-9	-8	0.5
start	1	0	0.5	1

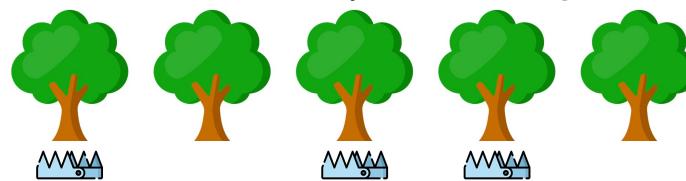
## Gridworld problem

**Feature:** grid cell features

**Prediction:** grid cell reward

**Action:** movement

Transition probability



## Snare-finding problem

**Feature:** location features

**Prediction:** snaring risk

**Action:** visit a location to remove snares

Transition probability



## Tuberculosis problem

**Feature:** patient features

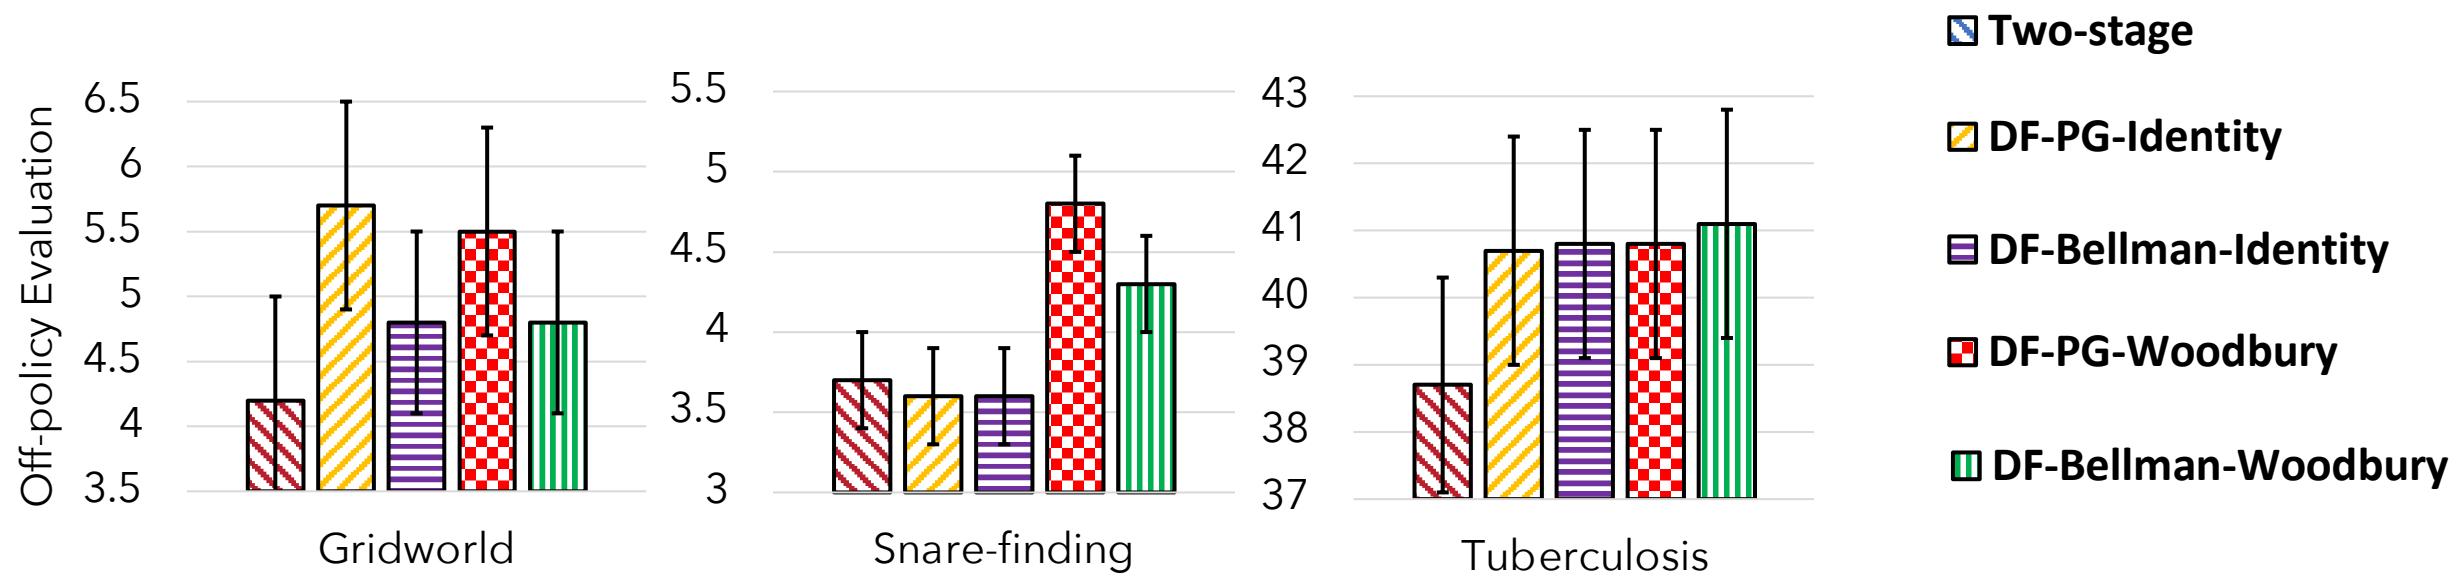
**Prediction:** patient transition probabilities

**Action:** visit a patient to intervene

# Methods

Standard two-stage approach		☒ Two-stage
Approximation	Optimality condition	
Identity	Policy gradient	Bellman error
Woodbury (low-rank)	<input checked="" type="checkbox"/> <b>DF-PG-Identity</b> <input checked="" type="checkbox"/> <b>DF-PG-Woodbury</b>	<input checked="" type="checkbox"/> <b>DF-Bellman-Identity</b> Futoma et al. (AISTATS 2020) <input checked="" type="checkbox"/> <b>DF-Bellman-Woodbury</b>

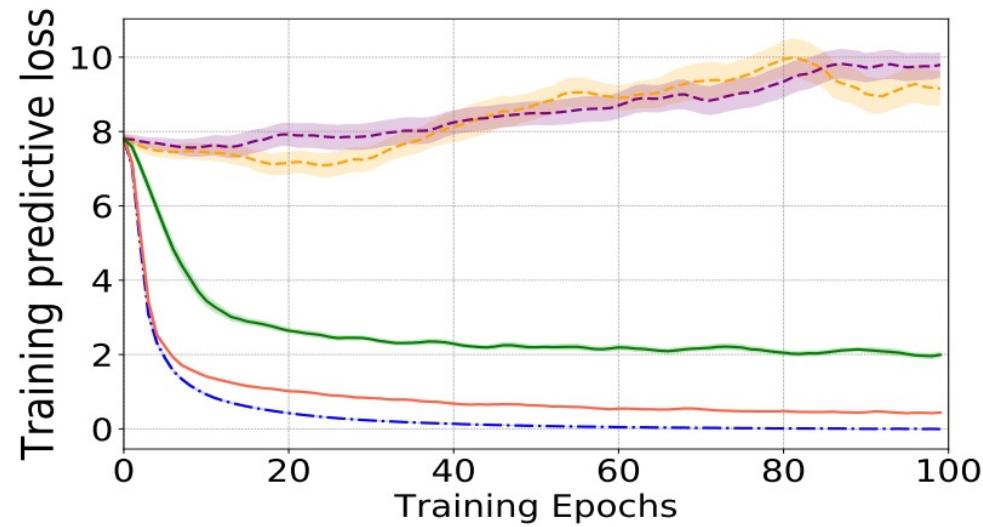
# Results



DFL outperforms two-stage  
Woodbury-based are more consistent than identity-based

# Analysis (snare-finding)

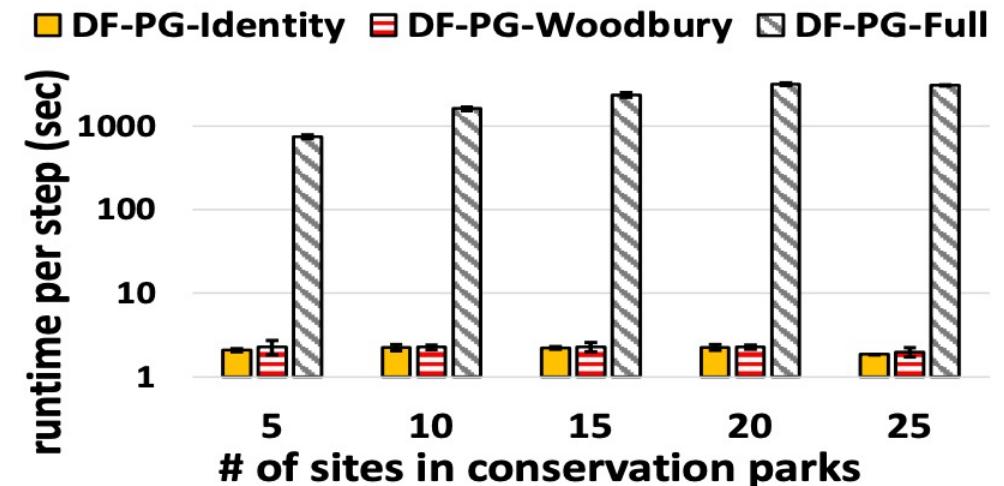
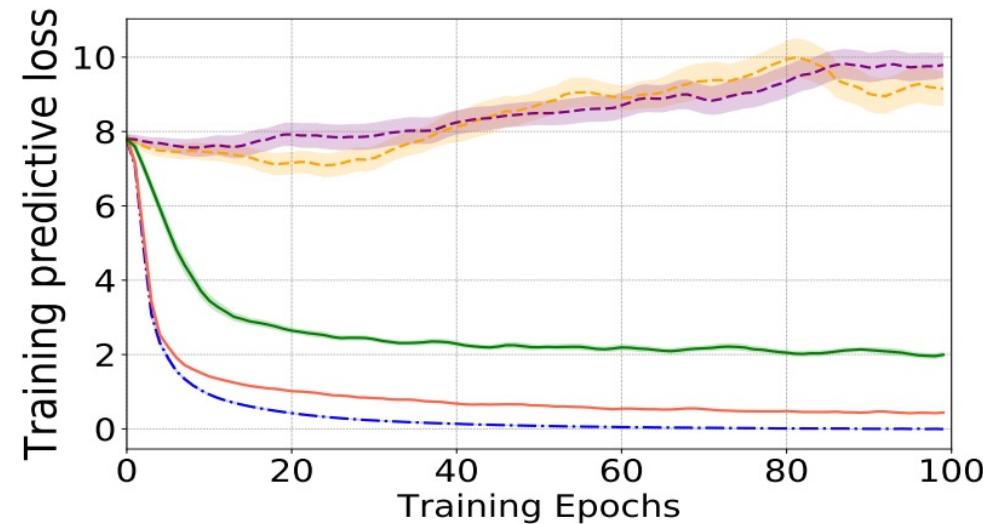
— TS    — DF-PG-Identity    -·- DF-Bellman-Identity    — DF-PG-Woodbury    — DF-Bellman-Woodbury



DFL models learn less accurate models than two-stage  
Hessian approximation saves ~3 orders of magnitude

# Analysis (snare-finding)

— TS    — DF-PG-Identity    -·- DF-Bellman-Identity    — DF-PG-Woodbury    — DF-Bellman-Woodbury



DFL models learn less accurate models than two-stage  
Hessian approximation saves ~3 orders of magnitude

# Why is DFL helpful?

Consider  $f(z, \theta) = \sum_i \theta_i z_i$ . Then,

$$\max_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} \sum_i \theta_i z_i = \max_{z \in Z} \sum_i \mathbb{E}[\theta_i | x] z_i$$

# Why is DFL helpful?

Consider  $f(z, \theta) = \sum_i \theta_i z_i$ .

$\max_{z \in Z} \sum_i m_{w^*}(x) z_i$  instead of  $\max_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} \sum_i \theta_i z_i$

s.t.  $w^* = \operatorname{argmin}_w \mathcal{L}(w, \text{training data})$

# Why is DFL helpful?

Consider  $f(z, \theta) = \sum_i \theta_i z_i$ .

$\max_z \sum_i m_{\omega^*}(x) z_i$  instead of  $\max_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} \sum_i \theta_i z_i$

s.t.  $w^* = \operatorname{argmin}_w \mathcal{L}(w, \text{training data})$

$z_i \geq 0, \sum_i z_i \leq 1$

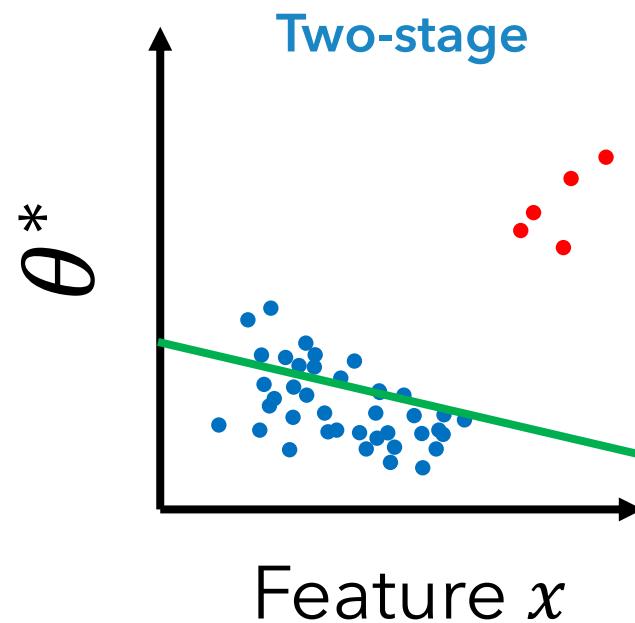
# Why is DFL helpful?

Consider  $f(z, \theta) = \sum_i \theta_i z_i$ .

$\max_z \sum_i m_{\omega^*}(x) z_i$  instead of  $\max_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} \sum_i \theta_i z_i$

s.t.  $w^* = \operatorname{argmin}_w \mathcal{L}(w, \text{training data})$

$$z_i \geq 0, \sum_i z_i \leq 1$$



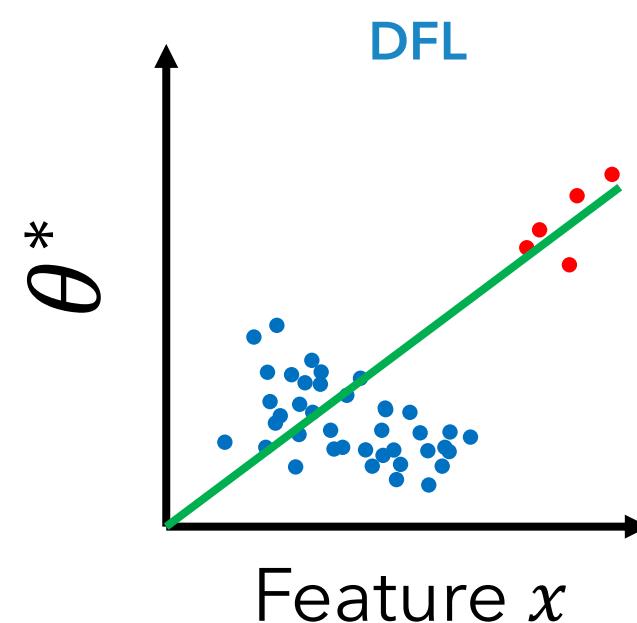
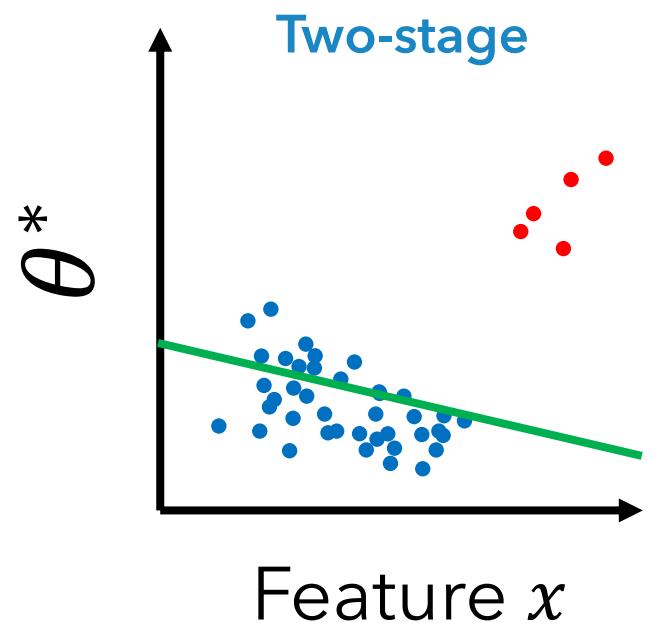
# Why is DFL helpful?

Consider  $f(z, \theta) = \sum_i \theta_i z_i$ .

$\max_z \sum_i m_{\omega^*}(x) z_i$  instead of  $\max_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} \sum_i \theta_i z_i$

s.t.  $w^* = \operatorname{argmin}_w \mathcal{L}(w, \text{training data})$

$$z_i \geq 0, \sum_i z_i \leq 1$$

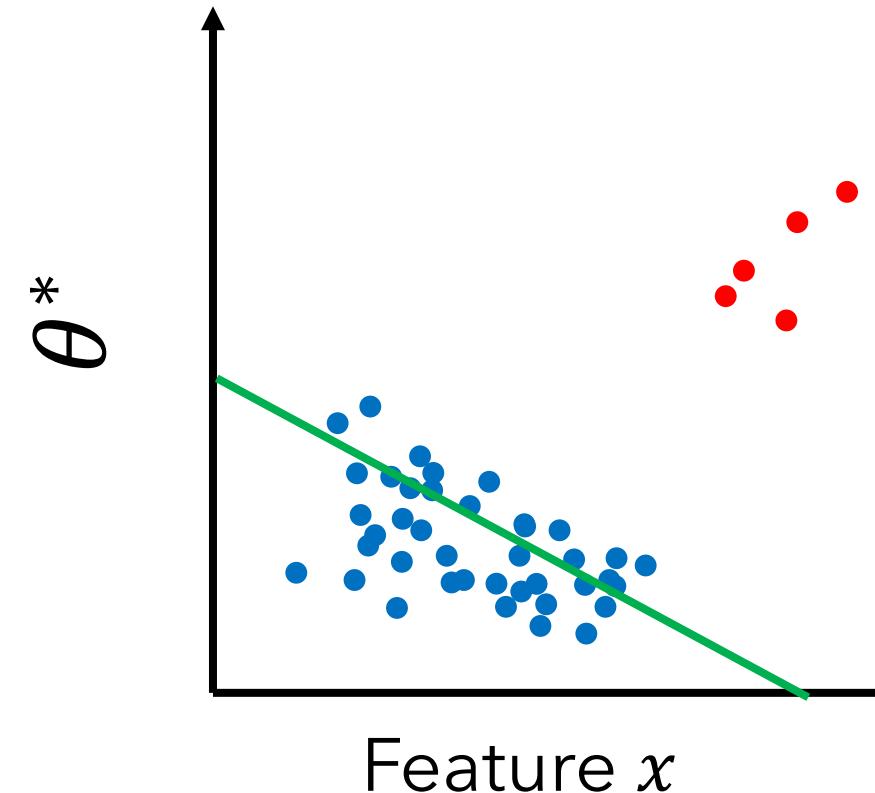
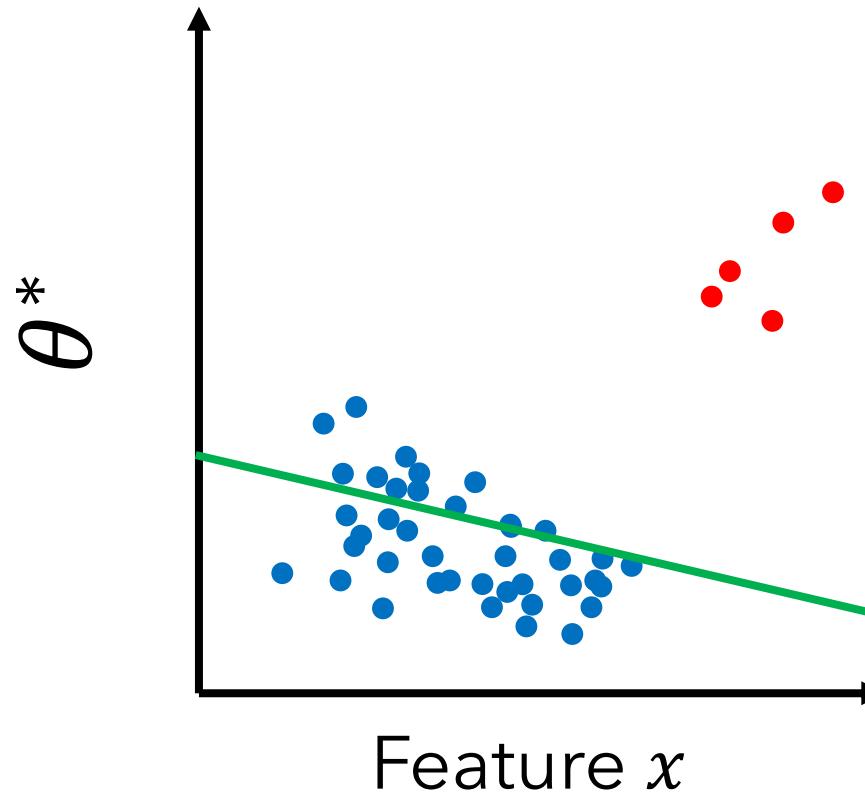


# The problem of zero gradient

$$\max_z \sum_i m_{\omega^*}(x) z_i$$

s.t.  $w^* = \operatorname{argmin}_w \mathcal{L}(w, \text{training data})$

$$z_i \geq 0, \sum_i z_i \leq 1$$

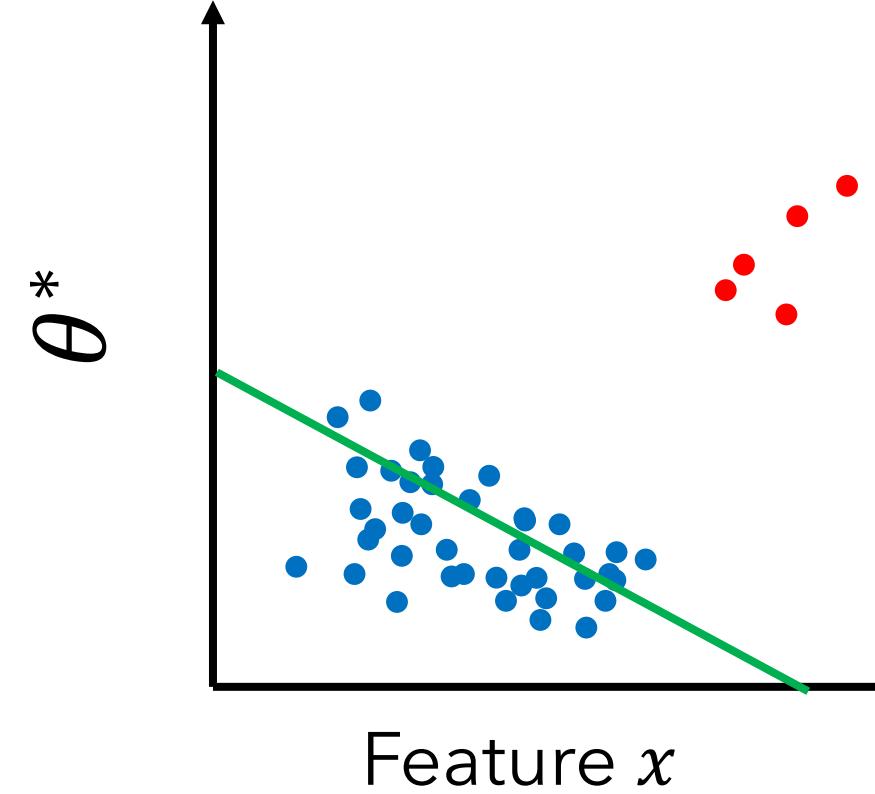
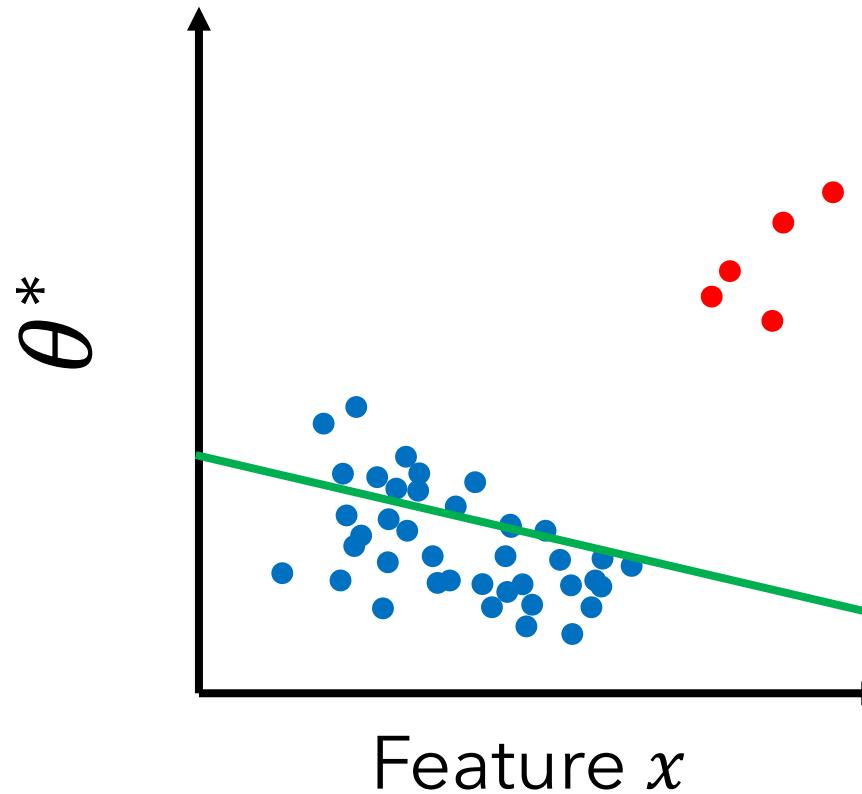


# The problem of zero gradient

$$\max_z \sum_i m_{\omega^*}(x) z_i + \lambda z_i^2$$

s.t.  $w^* = \operatorname{argmin}_w \mathcal{L}(w, \text{training data})$

$$z_i \geq 0, \sum_i z_i \leq 1$$

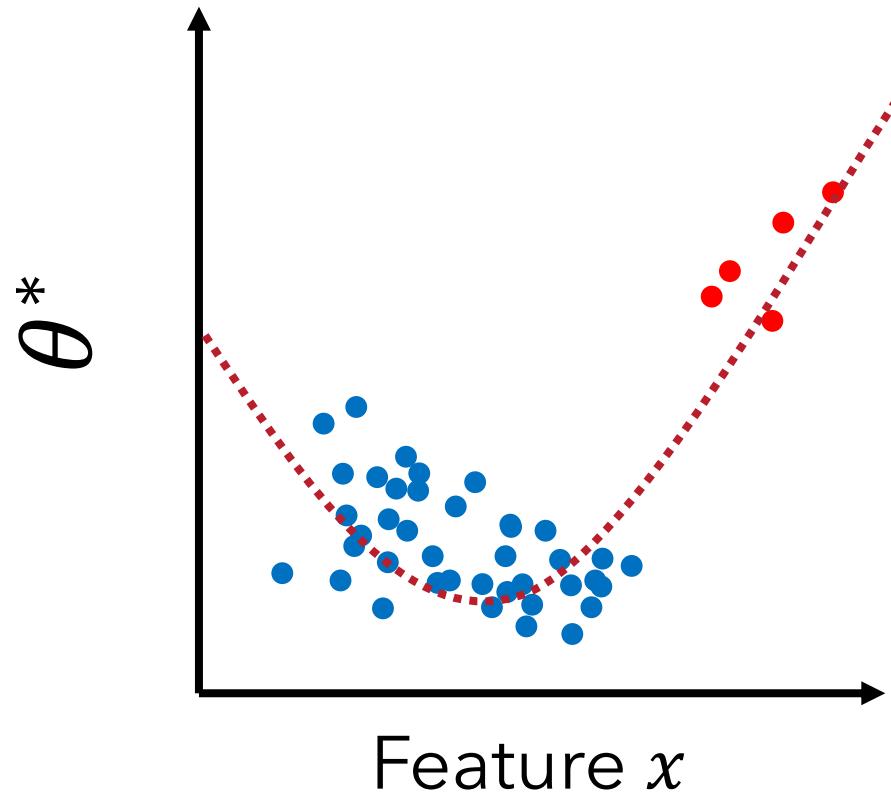


# The problem of non-convexity

$$\max_z \sum_i m_{\omega^*}(x) z_i + \lambda z_i^2$$

s.t.  $w^* = \operatorname{argmin}_w \mathcal{L}(w, \text{training data})$

$$z_i \geq 0, \sum_i z_i \leq 1$$

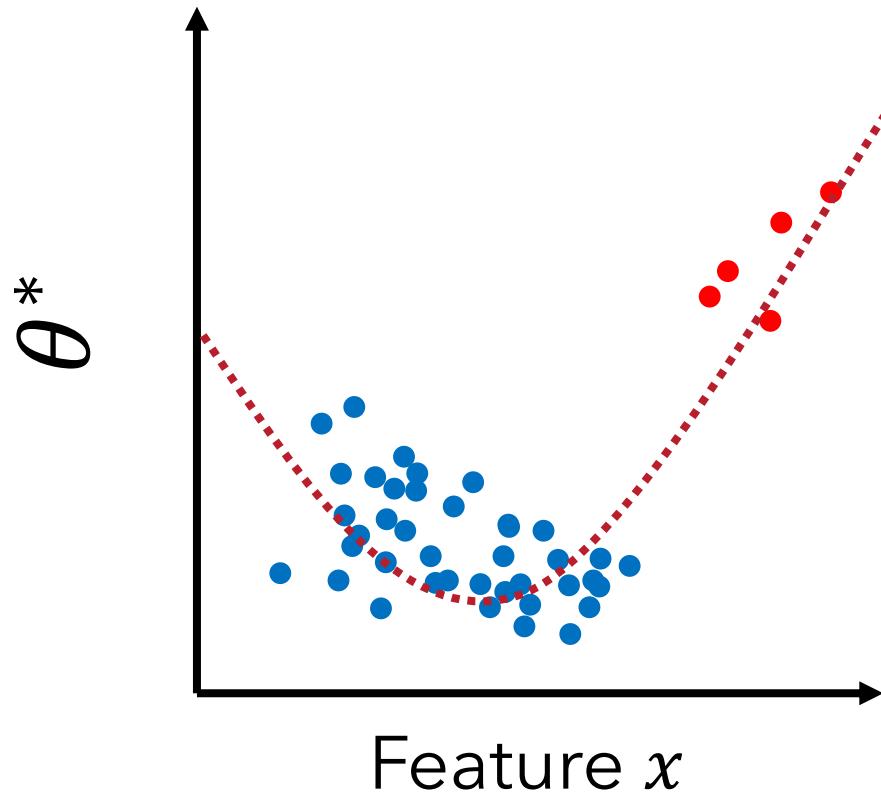


# The problem of non-convexity

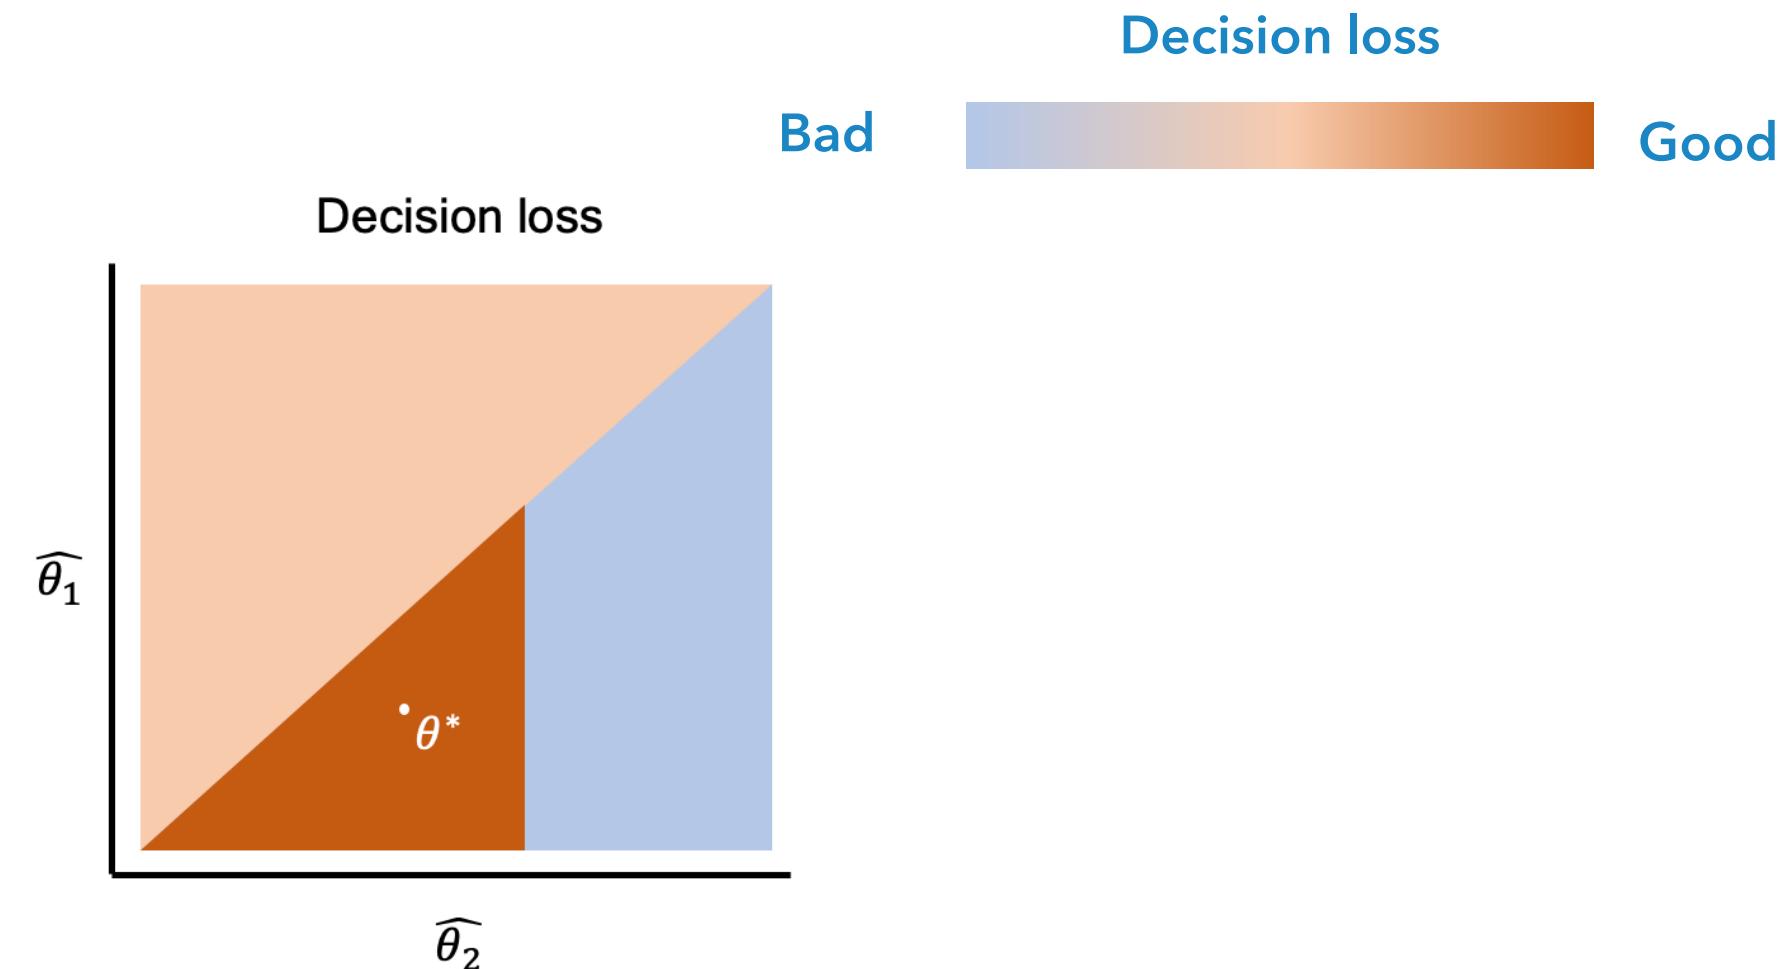
$$\max_{\mathbf{z}} \sum_i m_{\omega^*}(x) z_i + \lambda z_i^2$$

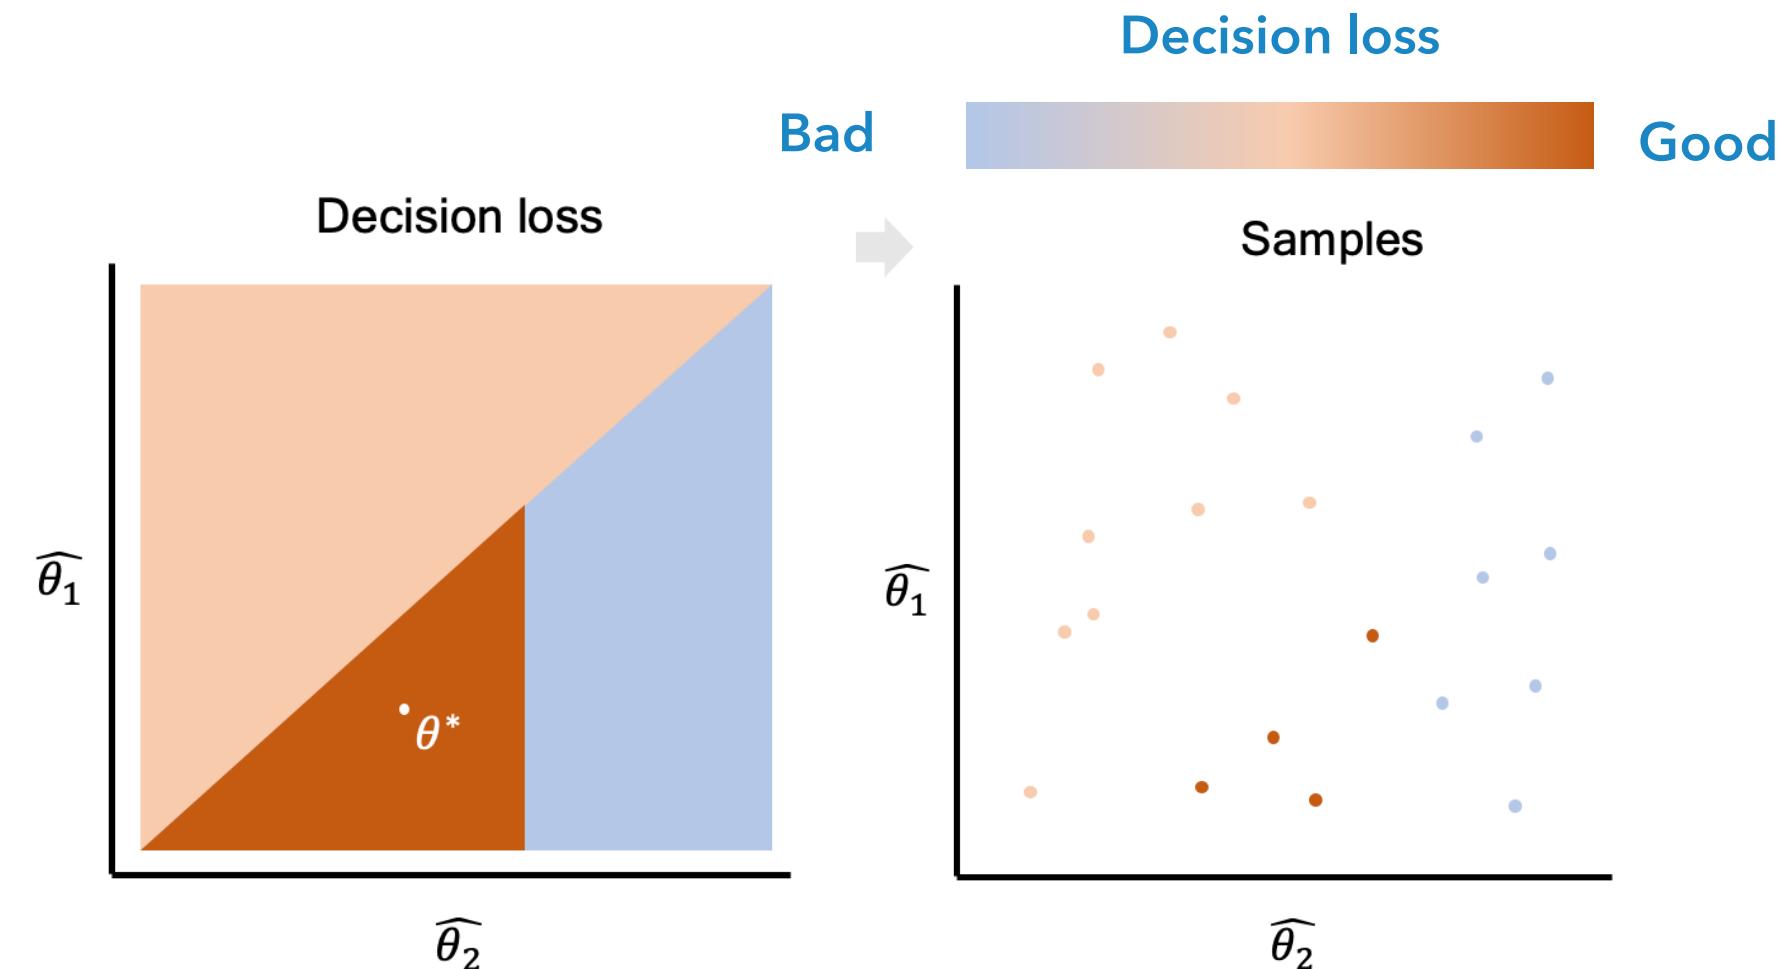
s.t.  $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \mathcal{L}(\mathbf{w}, \text{training data})$

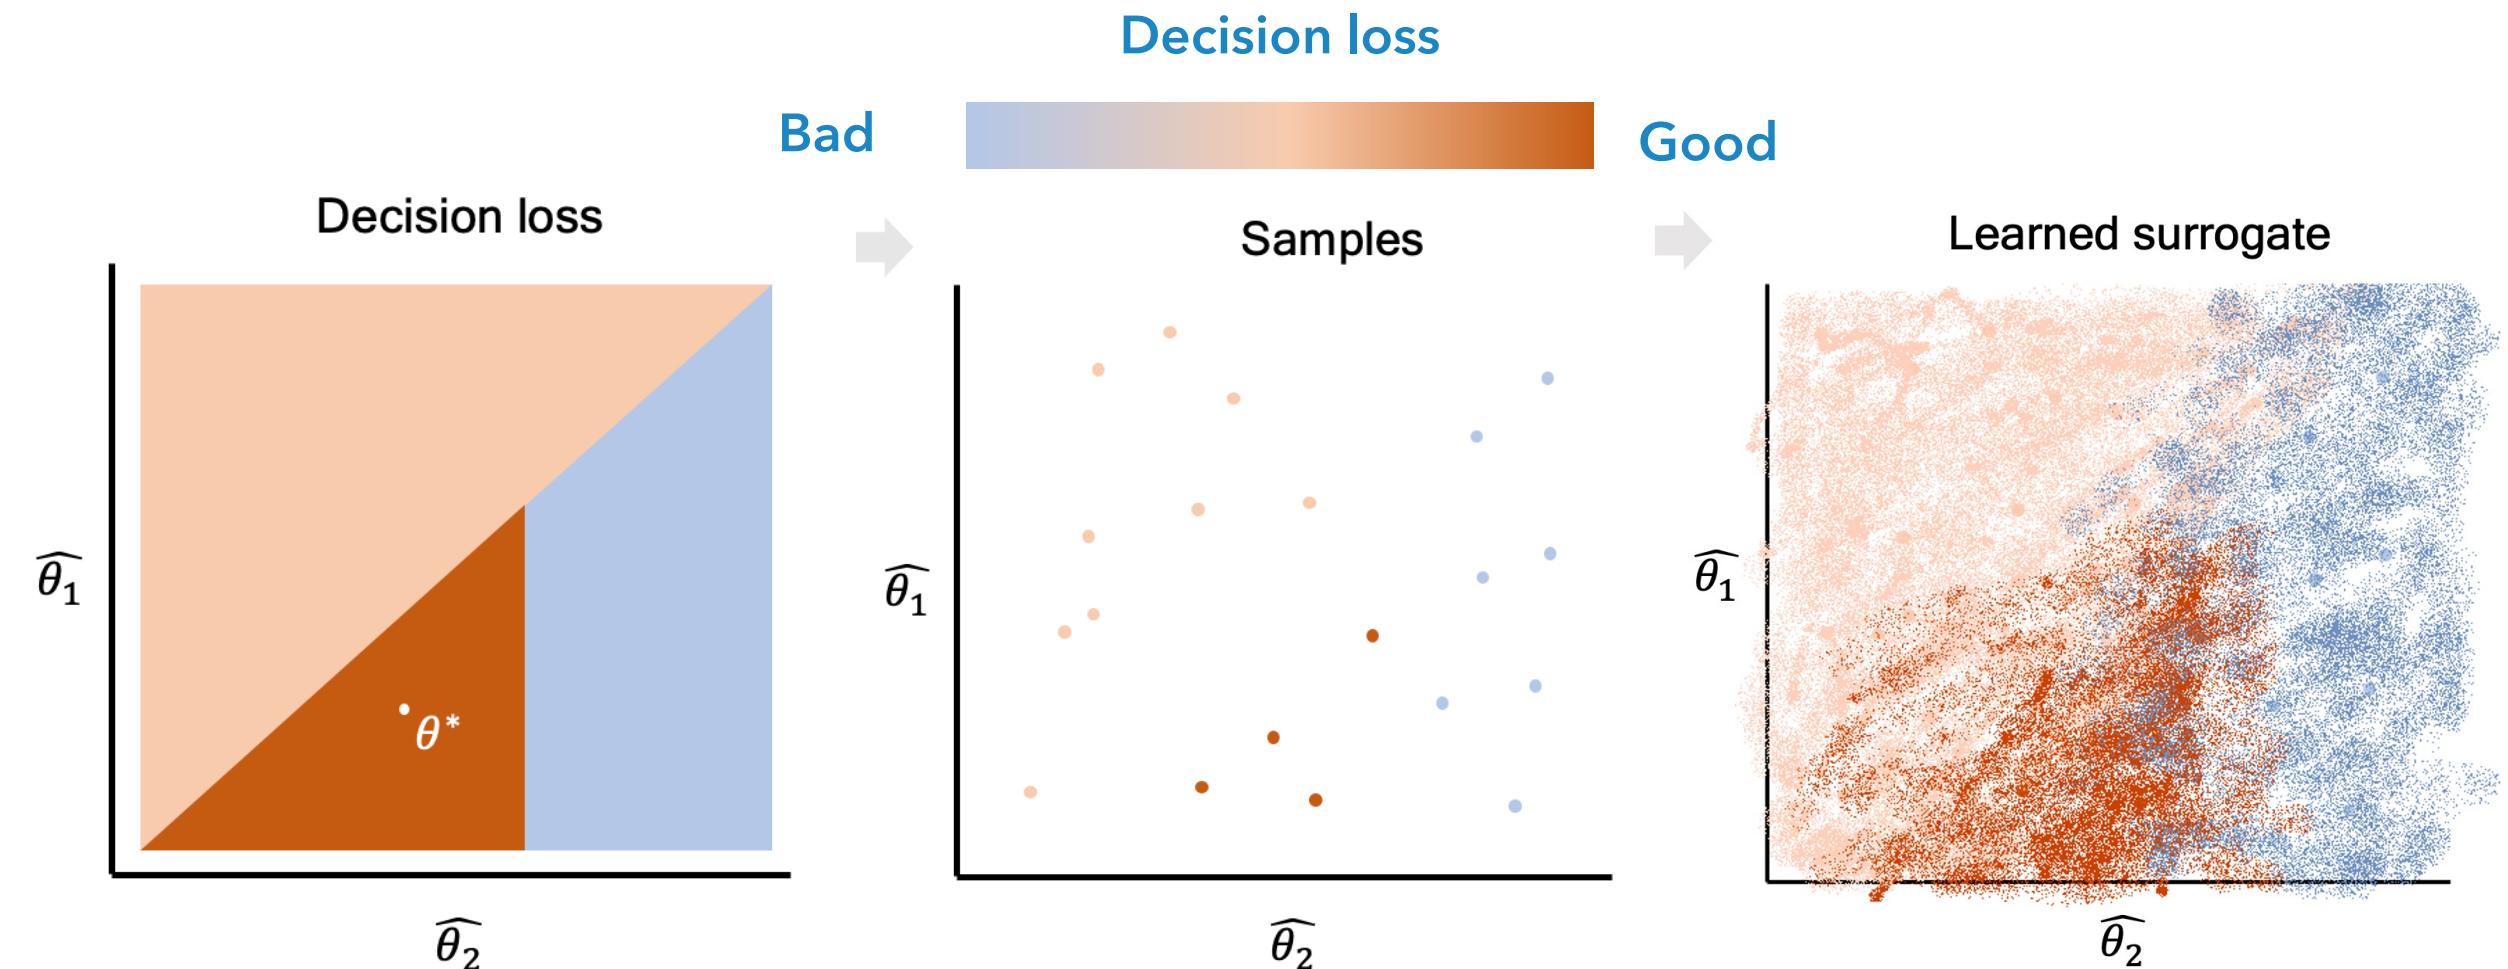
$$z_i \geq 0, \sum_i z_i \leq 1$$



Cool theory for linear objectives:  
Elmachtoub and Grigas (2022)







# Questions

1. Where to sample from the decision loss function?
2. Where is the local minimum?
3. What function to fit?



# Notation

Predicted environment parameter  $\hat{\theta}$

Decision induced by  $\hat{\theta}$ :  
 $z^*(\hat{\theta}) = \operatorname{argmin}_{z \in Z} f(z, \hat{\theta})$

Decision loss of  $\hat{\theta}$ :  
 $\ell(\hat{\theta}) = f(z^*(\hat{\theta}), \theta^*)$

## Standard DFL notation

Observable feature/context vector  $x \in X$

Unobserved ground truth environment parameter  $\theta^* \in \Theta$

$x$  and  $\theta^*$  drawn jointly from fixed distribution  $\mathcal{D}$

Decision/control variable  $z \in Z$

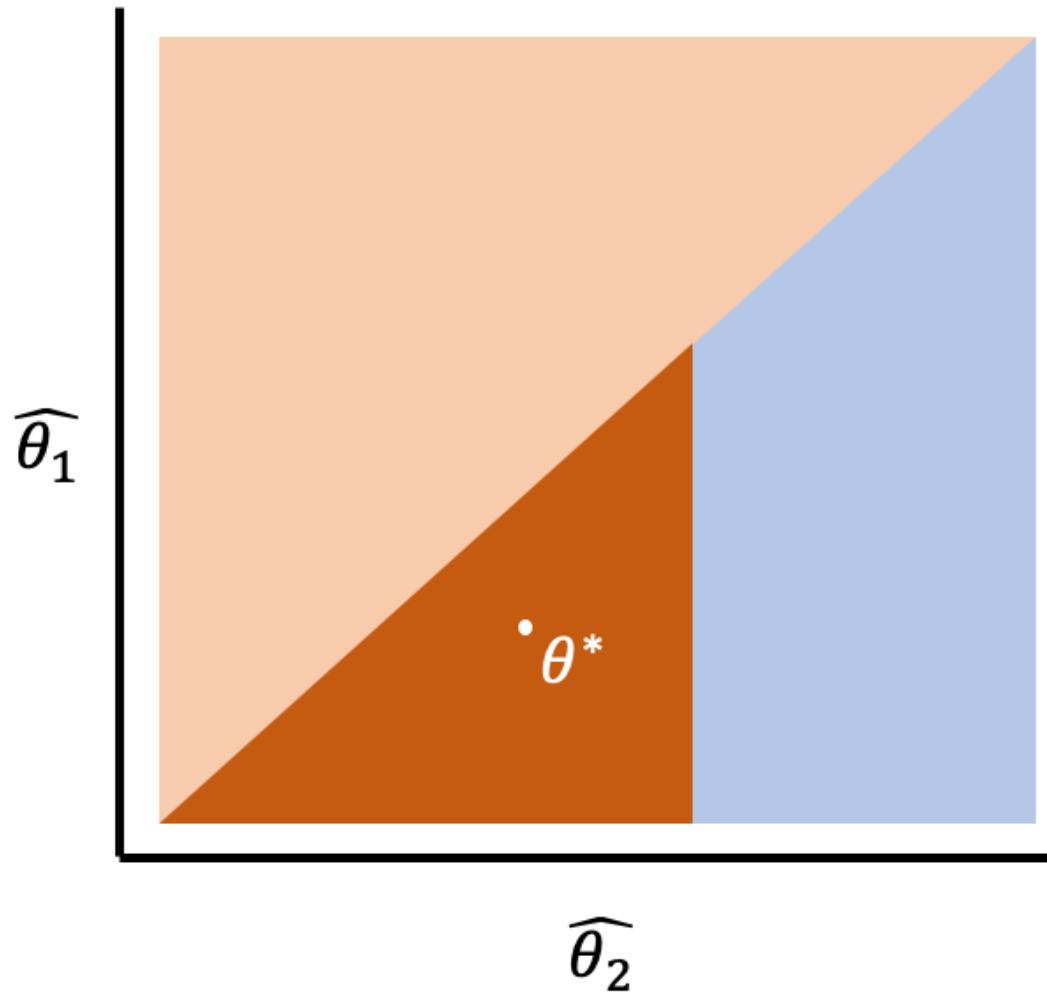
Objective function  $f(z, \theta^*)$

Goal:  $\min_{z \in Z} \mathbb{E}_{\theta \sim P(\Theta|x)} f(z, \theta)$

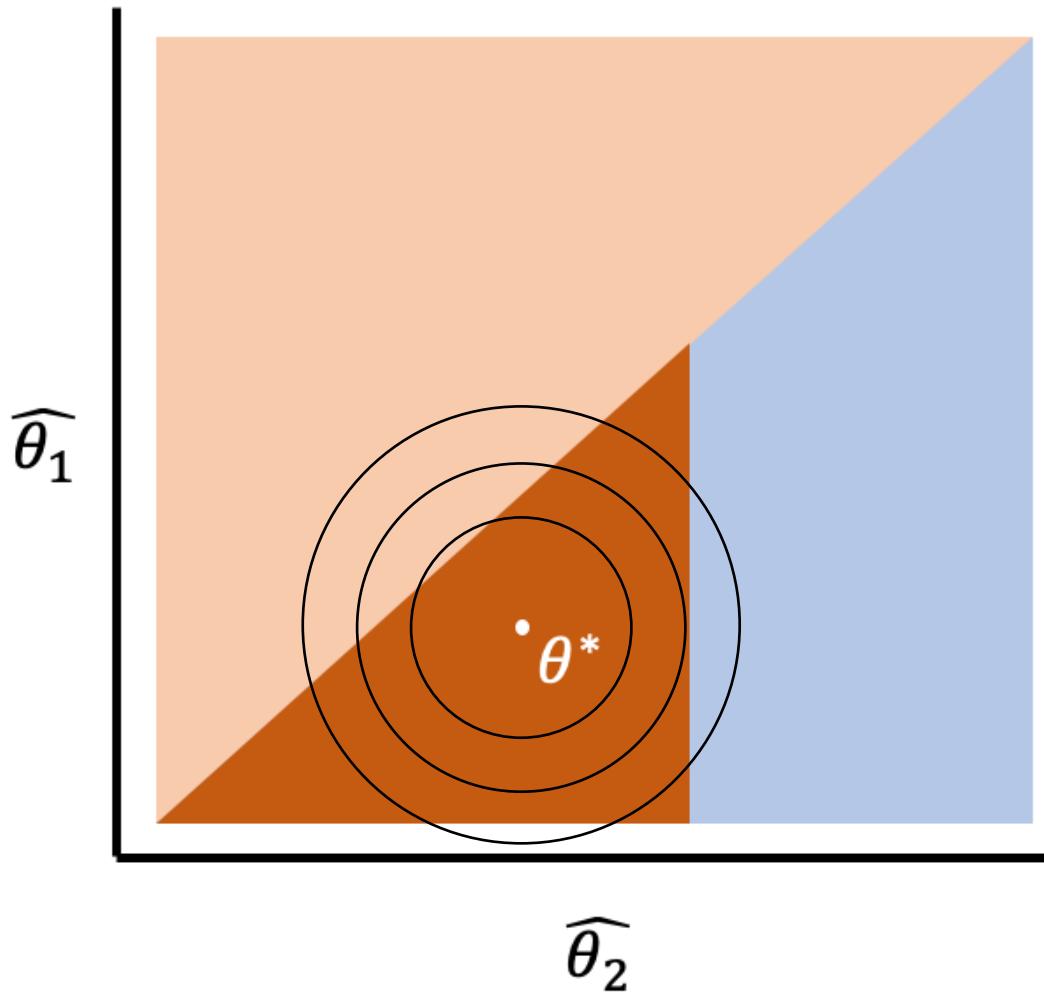


1. Sample a dataset of pairs  $\mathcal{D} = \{(\theta_j, \ell(\theta_j))\}$
2. Fit a parametric loss  $\hat{\ell}_\omega(\theta)$  to minimize MSE on training data
3. Train predictive model  $m: X \rightarrow \Theta$  using  $\hat{\ell}_\omega(\theta)$

## Decision loss

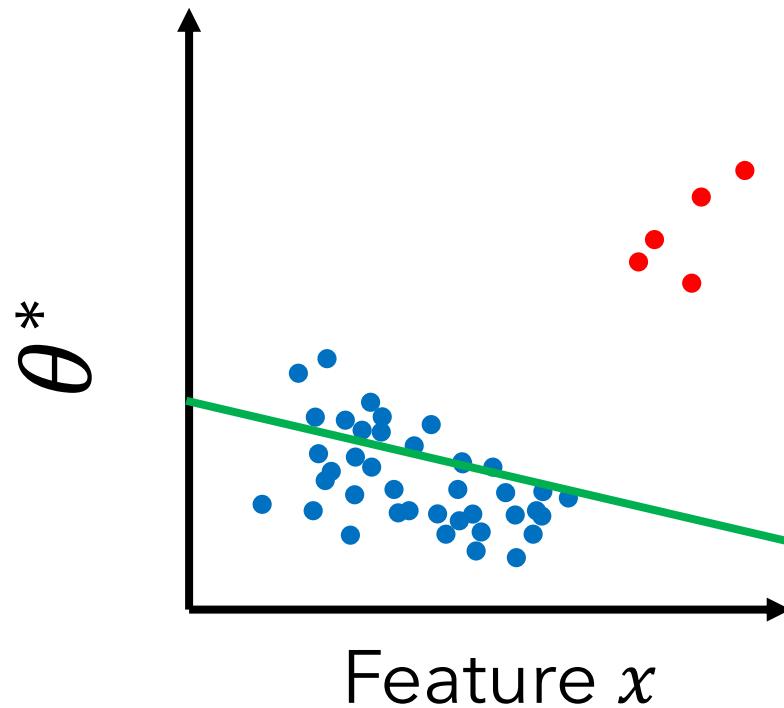


## Decision loss



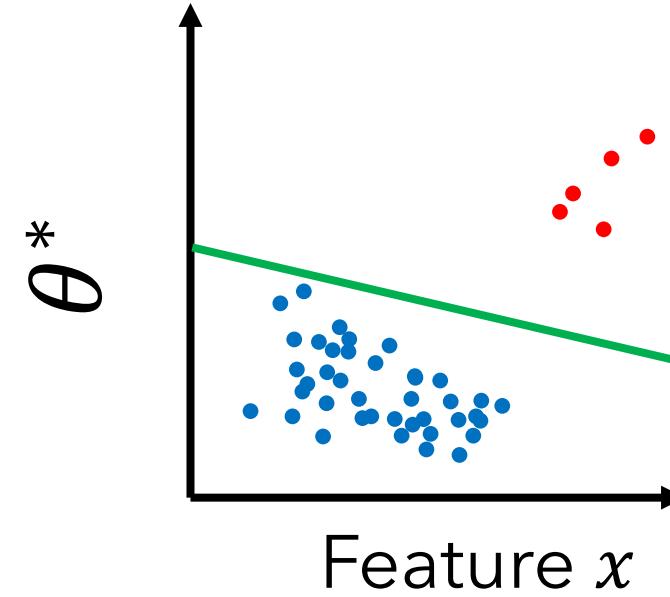
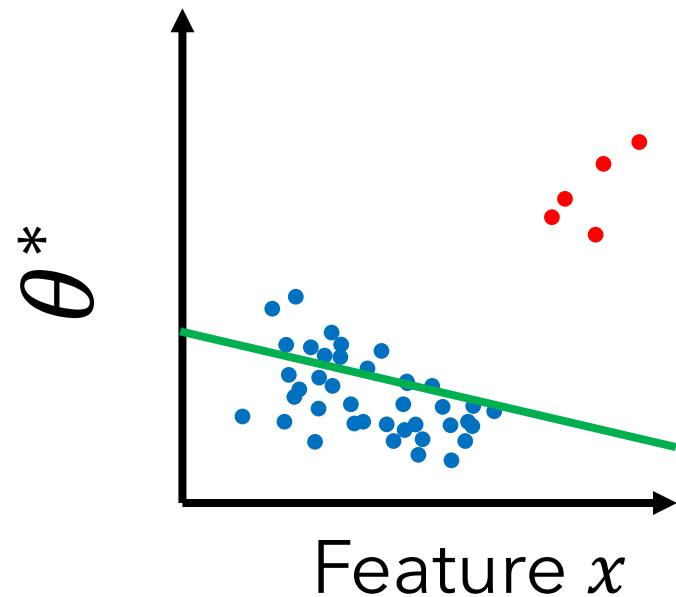
# Choosing a parametric family

1. Capture differing importance of dimensions of  $\theta$



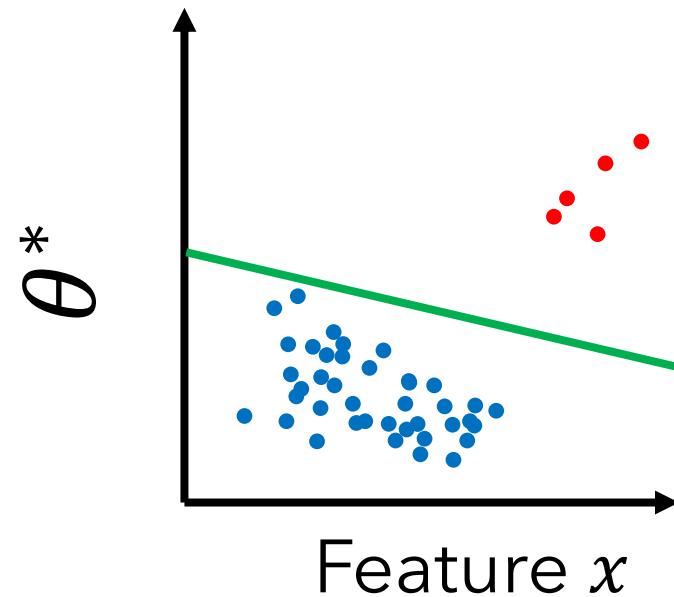
# Choosing a parametric family

1. Capture differing importance of dimensions of  $\theta$
2. Capture pairwise interactions between dimensions of  $\theta$



# Choosing a parametric family

1. Capture differing importance of dimensions of  $\theta$
2. Capture pairwise interactions between dimensions of  $\theta$
3. Capture asymmetry in over- vs. under-predicting



# Choosing a parametric family

Starting point: quadratic loss

$$(\hat{\theta} - \theta^*)^T H (\hat{\theta} - \theta^*)$$

where  $H$  is the learned parameter.

Low-rank, PSD decomposition:  $H = LL^T$



# Choosing a parametric family

Starting point: quadratic loss

$$(\hat{\theta} - \theta^*)^T H (\hat{\theta} - \theta^*)$$

where  $H$  is the learned parameter.

Low-rank, PSD decomposition:  $H = LL^T$

- ✓ Capture differing importance of dimensions of  $\theta$
- ✓ Capture pairwise interactions between entries of  $\theta$
- ✗ Capture asymmetry in over vs under predicting



# Choosing a parametric family

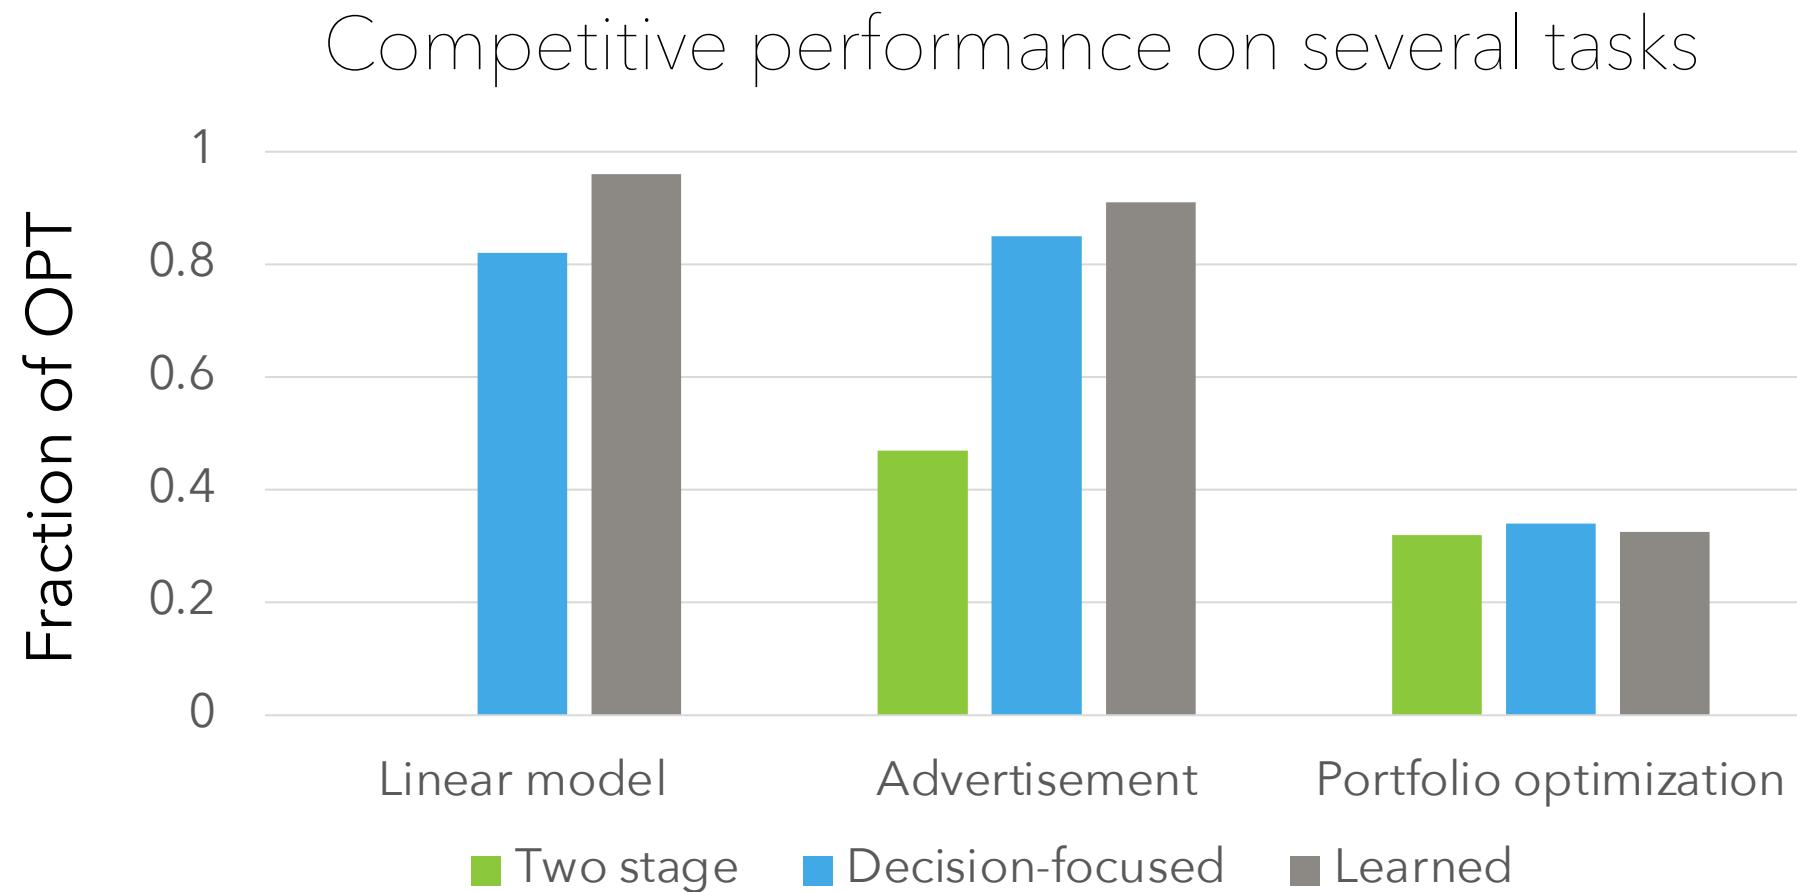
Refinement: *directed quadratic loss*.

Each  $i, j$  pair of coordinates gets a loss which depends on direction of errors for the prediction of each

Instantiate via four copies of the parameter  $L$  ( $H = LL^T$ )

$$\begin{aligned} L_{ij}^{++}, & \quad \text{if } \hat{y}_i - y_i \geq 0 \text{ and } \hat{y}_j - y_j \geq 0 \\ L_{ij}^{+-}, & \quad \text{if } \hat{y}_i - y_i \geq 0 \text{ and } \hat{y}_j - y_j < 0 \\ L_{ij}^{-+}, & \quad \text{if } \hat{y}_i - y_i < 0 \text{ and } \hat{y}_j - y_j \geq 0 \\ L_{ij}^{--}, & \quad \text{otherwise} \end{aligned}$$

# Experiments



# Conclusion/future directions

1. Cost of training
2. Exploiting correlation structure
3. Accessibility of code/methods
4. Benchmarks



# References

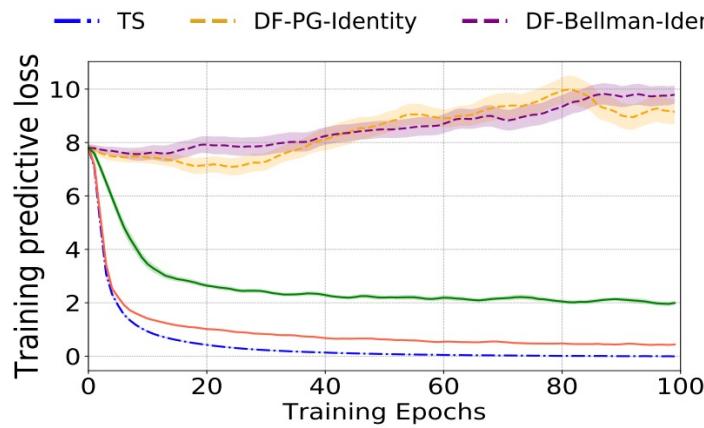
- Wilder et al. Melding the data-decisions pipeline: decision-focused learning for combinatorial optimization. AAAI 2019.
- Cameron et al. The Perils of Learning Before Optimizing. AAAI 2022.
- Agrawal et al. Price of correlations in stochastic optimization. Operations Research, 2012.
- Wang et al. Learning MDPs from Features: Predict-Then-Optimize for Sequential Decision Problems by Reinforcement Learning. NeurIPS 2021.**
- Mate et al. Collapsing Bandits and Their Application to Public Health Interventions. NeurIPS 2020.
- Futoma et al. POPCORN: Partially Observed Prediction COnstrained ReiNforcement learning. AISTATS 2020.
- Elmachtoub and Grigas. Smart "Predict, then Optimize". Management Science, 2022.
- Shah et al. Decision-Focused Learning without Decision-Making: Learning Locally Optimized Decision Losses. arXiv, 2022.**

## References (tutorials)

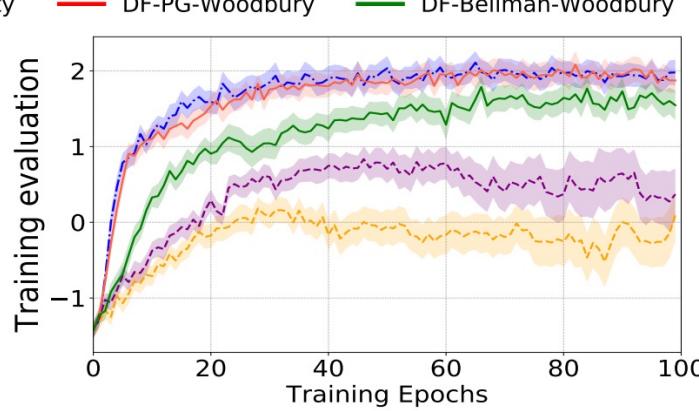
Deep implicit layers (NeurIPS 2020): <http://implicit-layers-tutorial.org/>

Deep declarative networks (ECCV 2020):  
<https://anucvml.github.io/ddn-eccv2020/>

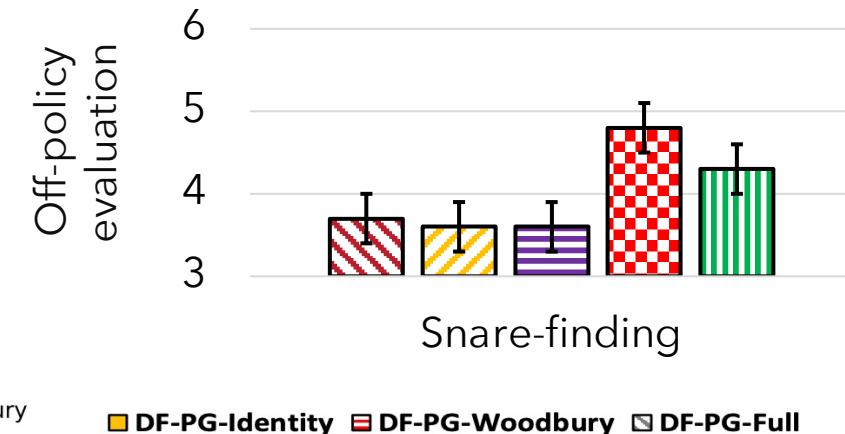
# Analysis (snare-finding)



(a) Predictive loss in training MDPs



(b) Performance in training MDPs



DFL models learn less accurate models than two-stage  
Two-stage scores better on train but worse on test

# Price of correlation

Definition (price of correlation): for an objective function  $f$ , the price of correlation is ratio  $\frac{f(x',y)}{f(x^*,y)}$ , where

- $x^* = \min_{x \in X} \mathbb{E}_{y \sim P(y|\theta)} f(x, y)$
- $x' = \min_{x \in X} \mathbb{E}_{y \sim P(y|\theta)} f(x, y)$  assuming  $y_i$  mutually independent